



UNIVERSIDAD  
**NACIONAL**  
DE COLOMBIA

# Medical Image Segmentation in a Multiple Labelers Context: Application to the Study of Histopathology

Brandon Lotero Londoño

Universidad Nacional de Colombia  
Faculty of Engineering and Architecture  
Department of Electric, Electronic and Computing Engineering  
Manizales, Colombia  
2025



# Medical Image Segmentation in a Multiple Labelers Context: Application to the Study of Histopathology

Brandon Lotero Londoño

Dissertation submitted in partial fulfillment to receive the grade of:  
Master in Engineering - Industrial Automation

Advisor:

Prof. Andrés Marino Álvarez-Meza, Ph.D.

Co-advisor:

Prof. Germán Castellanos-Domínguez, Ph.D.

Academic research group:

Signal Processing and Recognition Group - SPRG

Universidad Nacional de Colombia

Faculty of Engineering and Architecture

Department of Electric, Electronic and Computing Engineering

Manizales, Colombia

2025





# **Segmentación de Imágenes Médicas en un Contexto de Múltiples Anotadores: Aplicación al Estudio de Histopatologías**

**Brandon Lotero Londoño**

Disertación presentada como requisito parcial para recibir el título de:

**Magíster en Ingeniería - Automatización Industrial**

Director:

Prof. Andrés Marino Álvarez-Meza, Ph.D.

Codirector:

Prof. Germán Castellanos-Domínguez, Ph.D.

Grupo de investigación:

Grupo de Control y Procesamiento Digital de Señales - GCPDS

Universidad Nacional de Colombia

Facultad de Ingeniería y Arquitectura

Departamento de Ingeniería Eléctrica, Electrónica y Computación

Manizales, Colombia

2025



## ACKNOWLEDGEMENTS

45 I would like to express my deepest gratitude to:

46 My beloved Camila, for her unwavering love, support, and patience throughout this  
47 journey.

48 My family, especially my parents, for their unconditional support during challenging  
49 times. I look forward to sharing many more joyful moments together.

50 My advisors, Prof. Andrés Marino Álvarez-Meza, Prof. Julián Gil, and Prof. Germán  
51 Castellanos-Domínguez, for their invaluable guidance and mentorship.

52 My friends and colleagues at GCPDS, particularly Lucas, Santiago, Marcos, and  
53 Rafael, for their support and camaraderie throughout this academic journey.

54 The Universidad Nacional de Colombia, for providing the resources and facilities  
55 that enabled me to pursue this Master's degree. I hope this work contributes  
56 meaningfully to the field of early cancer detection.



## ABSTRACT

60 Breast cancer remains one of the deadliest diseases affecting women worldwide,  
61 making accurate diagnosis and treatment planning crucial for improving patient  
62 outcomes. Semantic segmentation of histopathology images plays a vital role in  
63 this process, as it enables precise identification and analysis of tissue structures.  
64 However, the current approaches face significant challenges in real-world  
65 scenarios where multiple annotators with varying expertise and reliability  
66 contribute to the labeling process. This crowdsourcing paradigm, while  
67 cost-effective, introduces inconsistencies that can compromise the quality of the  
68 segmentation results.

69 This thesis addresses two fundamental challenges in medical image segmentation:  
70 the inherent uncertainty in multiple annotator scenarios and the need for robust  
71 segmentation models that can handle varying image quality and annotator  
72 expertise. These challenges arise from the complex nature of histopathology  
73 images, where tissue structures can be ambiguous, and the varying levels of  
74 expertise among annotators lead to inconsistent labeling patterns. Additionally,  
75 the high cost and time requirements for expert annotations create a need for  
76 efficient solutions that can work with crowdsourced data.

77 To address these challenges, this thesis proposes a comprehensive approach that  
78 combines novel loss functions with robust deep learning architectures to handle  
79 multiple annotator scenarios effectively.

80 The first objective introduces a novel loss function that implicitly infers optimal  
81 segmentation while assessing labeler reliability without requiring explicit  
82 performance inputs. This approach differs from conventional methods by  
83 eliminating the need for direct supervision of labeler performance, instead using  
84 an intermediate reliability map that enables the model to prioritize information  
85 from reliable labelers while disregarding noisy labels. The mathematical  
86 foundation of this approach lies in the probabilistic modeling of annotator  
87 reliability across the input space.

88 The second objective focuses on the development of a robust segmentation  
89 architecture that combines the strengths of U-shaped deep learning models with  
90 the proposed reliability-aware loss function. This combination enables the model  
91 to learn both the segmentation task and the reliability patterns simultaneously,  
92 leading to more accurate and consistent results. The architecture's key innovation  
93 lies in its ability to adapt to varying image quality and annotator expertise levels  
94 without requiring explicit performance metrics.

95 The third objective involves validating the proposed approach through extensive  
96 experimentation on both synthetic datasets and real-world histopathology images.  
97 The evaluation demonstrates superior performance compared to state-of-the-art  
98 approaches in terms of segmentation accuracy and uncertainty quantification. The  
99 results show that the proposed method can effectively handle inconsistent  
100 annotations while maintaining high segmentation accuracy, making it particularly  
101 valuable for clinical applications.

102 The research makes significant contributions to the field by providing a robust  
103 solution for handling inconsistent annotations in medical image segmentation,  
104 particularly valuable in histopathology where expert annotations are expensive  
105 and time-consuming to obtain. The proposed approach shows promising results in  
106 reducing annotation costs while maintaining high segmentation accuracy, making  
107 it particularly relevant for clinical applications where precise segmentation of  
108 tissue structures is crucial for diagnosis and treatment planning. This work opens

109 new research directions in the areas of crowdsourced medical image analysis,  
110 uncertainty quantification in deep learning, and the development of more efficient  
111 annotation protocols for medical imaging.

112 **Keywords:** Crowdsourcing, Multiple Annotators, Histopathology, Breast Cancer,  
113 Semantic Image Segmentation, Gaussian Processes, Deep Learning, Medical Image  
114 Analysis





116 El cáncer de mama sigue siendo una de las enfermedades más mortales en la  
117 población femenina en todo el mundo, lo que hace que el diagnóstico preciso y la  
118 planificación del tratamiento sean cruciales para mejorar las condiciones de los  
119 pacientes. La segmentación semántica de imágenes histopatológicas juega un  
120 papel vital en este proceso, ya que permite la identificación y análisis preciso de  
121 estructuras tisulares. Sin embargo, los enfoques actuales enfrentan desafíos  
122 significativos en escenarios del mundo real donde múltiples anotadores con  
123 diferentes niveles de experiencia y confiabilidad contribuyen al proceso de  
124 etiquetado. Este paradigma de crowdsourcing, aunque rentable, introduce  
125 inconsistencias que pueden comprometer la calidad de los resultados de  
126 segmentación.

127 Esta tesis aborda dos desafíos fundamentales en la segmentación de imágenes  
128 médicas: la incertidumbre inherente en escenarios de múltiples anotadores y la  
129 necesidad de modelos de segmentación robustos que puedan manejar la calidad  
130 variable de las imágenes y la experiencia de los anotadores. Estos desafíos surgen  
131 de la naturaleza compleja de las imágenes histopatológicas, donde las estructuras  
132 tisulares pueden ser ambiguas, y los diferentes niveles de experiencia entre los  
133 anotadores conducen a patrones de etiquetado inconsistentes. Además, el alto  
134 costo y los requisitos de tiempo para las anotaciones de expertos crean la  
135 necesidad de soluciones eficientes que puedan trabajar con datos crowdsourced.

El primer objetivo introduce una función de pérdida novedosa que infiere implícitamente la segmentación óptima mientras evalúa la confiabilidad de los etiquetadores sin requerir entradas explícitas sobre su rendimiento. Este enfoque difiere de los métodos convencionales al eliminar la necesidad de supervisión directa del rendimiento de los etiquetadores, utilizando en su lugar un mapa de confiabilidad intermedio que permite al modelo priorizar la información de etiquetadores confiables mientras descarta etiquetas ruidosas. La base matemática de este enfoque radica en el modelado probabilístico de la confiabilidad de los anotadores a través del espacio de entrada.

El segundo objetivo se centra en el desarrollo de una arquitectura de segmentación robusta que combina las fortalezas de los modelos de aprendizaje profundo en forma de U con la función de pérdida propuesta que tiene en cuenta la confiabilidad. Esta combinación permite que el modelo aprenda tanto la tarea de segmentación como los patrones de confiabilidad simultáneamente, lo que lleva a resultados más precisos y consistentes. La innovación clave de la arquitectura radica en su capacidad para adaptarse a diferentes niveles de calidad de imagen y experiencia de los anotadores sin requerir métricas de rendimiento explícitas.

El tercer objetivo valida el enfoque propuesto a través de experimentación exhaustiva tanto en conjuntos de datos sintéticos como en imágenes histopatológicas del mundo real. La evaluación demuestra un rendimiento superior en comparación con los enfoques más avanzados en términos de precisión de segmentación y cuantificación de incertidumbre. Los resultados muestran que el método propuesto puede manejar efectivamente anotaciones inconsistentes mientras mantiene una alta precisión de segmentación, lo que lo hace particularmente valioso para aplicaciones clínicas.

La investigación hace contribuciones significativas al campo al proporcionar una solución robusta para manejar anotaciones inconsistentes en la segmentación de imágenes médicas, particularmente valiosa en histopatología donde las anotaciones de expertos son costosas y requieren mucho tiempo para obtenerlas.

165 El enfoque propuesto muestra resultados prometedores en la reducción de costos  
166 de anotación mientras mantiene una alta precisión de segmentación, lo que lo hace  
167 particularmente relevante para aplicaciones clínicas donde la segmentación precisa  
168 de estructuras tisulares es crucial para el diagnóstico y la planificación del  
169 tratamiento. Este trabajo abre nuevas direcciones de investigación en las áreas de  
170 análisis de imágenes médicas crowdsourced, cuantificación de incertidumbre en  
171 aprendizaje profundo y el desarrollo de protocolos de anotación más eficientes  
172 para imágenes médicas.

173 **Palabras clave:** Crowdsourcing, Múltiples Anotadores, Histopatología, Cáncer de  
174 Mama, Segmentación Semántica de Imágenes, Procesos Gaussianos, Aprendizaje  
175 Profundo, Análisis de Imágenes Médicas



|     |                                                                 |              |
|-----|-----------------------------------------------------------------|--------------|
| 177 | <b>Acknowledgements</b>                                         | <b>vii</b>   |
| 178 | <b>Abstract</b>                                                 | <b>ix</b>    |
| 179 | <b>Resumen</b>                                                  | <b>xiii</b>  |
| 180 | <b>Contents</b>                                                 | <b>xx</b>    |
| 181 | <b>List of figures</b>                                          | <b>xxiv</b>  |
| 182 | <b>List of tables</b>                                           | <b>xxv</b>   |
| 183 | <b>List of algorithms</b>                                       | <b>xxvii</b> |
| 184 | <b>Abbreviations</b>                                            | <b>xxix</b>  |
| 185 | <b>1 Introduction</b>                                           | <b>1</b>     |
| 186 | 1.1 Motivation . . . . .                                        | 1            |
| 187 | 1.2 Problem Statement . . . . .                                 | 8            |
| 188 | 1.2.1 Variability in Expertise Levels . . . . .                 | 9            |
| 189 | 1.2.2 Technical Constraints and Image Quality . . . . .         | 9            |
| 190 | 1.2.3 Research Question . . . . .                               | 10           |
| 191 | 1.3 State of the Art . . . . .                                  | 10           |
| 192 | 1.3.1 Facing Annotation Variability in Medical images . . . . . | 12           |

|     |          |                                                           |           |
|-----|----------|-----------------------------------------------------------|-----------|
| 193 | 1.3.2    | Facing Noisy Annotations and Low-Quality Data . . . . .   | 20        |
| 194 | 1.4      | Aims . . . . .                                            | 21        |
| 195 | 1.4.1    | General Aim . . . . .                                     | 21        |
| 196 | 1.4.2    | Specific Aims . . . . .                                   | 23        |
| 197 | <b>2</b> | <b>Conceptual preliminaries</b>                           | <b>25</b> |
| 198 | 2.1      | Modern Concept of Digital Image . . . . .                 | 25        |
| 199 | 2.1.1    | Types of Digital Images . . . . .                         | 25        |
| 200 | 2.1.2    | Mathematical Representations . . . . .                    | 26        |
| 201 | 2.2      | Digital Histopathological Images . . . . .                | 29        |
| 202 | 2.2.1    | Whole Slide Imaging (WSI) . . . . .                       | 29        |
| 203 | 2.2.2    | Regions of Interest (ROI) . . . . .                       | 31        |
| 204 | 2.2.3    | Staining Techniques . . . . .                             | 33        |
| 205 | 2.3      | Deep Learning Fundamentals . . . . .                      | 34        |
| 206 | 2.3.1    | Learning Paradigms . . . . .                              | 34        |
| 207 | 2.3.2    | Architecture and Training . . . . .                       | 35        |
| 208 | 2.3.3    | Challenges and Solutions . . . . .                        | 36        |
| 209 | 2.3.4    | Deep Learning Frameworks . . . . .                        | 36        |
| 210 | 2.4      | Datasets and Data Sources . . . . .                       | 37        |
| 211 | 2.4.1    | Oxford-IIIT Pet Dataset . . . . .                         | 38        |
| 212 | 2.4.2    | Real histopathology datasets . . . . .                    | 39        |
| 213 | <b>3</b> | <b>Chained Gaussian Processes</b>                         | <b>43</b> |
| 214 | 3.1      | Background and Related Methods . . . . .                  | 43        |
| 215 | 3.1.1    | Kernel Alignment-Based Annotator Relevance Analysis . . . | 43        |
| 216 | 3.1.2    | Localized Kernel Alignment-Based Annotator Relevance      |           |
| 217 |          | Analysis . . . . .                                        | 44        |
| 218 | 3.1.3    | Chained Gaussian Processes Model . . . . .                | 45        |
| 219 | 3.1.4    | Regularized Chained Deep Neural Network . . . . .         | 46        |
| 220 | 3.2      | Tooling and Implementation . . . . .                      | 47        |
| 221 | 3.2.1    | GPflow Overview . . . . .                                 | 48        |
| 222 | 3.2.2    | GPflux Framework . . . . .                                | 49        |

|     |          |                                                                                     |           |
|-----|----------|-------------------------------------------------------------------------------------|-----------|
| 223 | 3.3      | Experimental Setup . . . . .                                                        | 50        |
| 224 | 3.3.1    | Classification . . . . .                                                            | 50        |
| 225 | 3.3.2    | Regression . . . . .                                                                | 52        |
| 226 | 3.4      | Results . . . . .                                                                   | 53        |
| 227 | 3.5      | Conclusion . . . . .                                                                | 53        |
| 228 | 3.5.1    | Comparative Analysis . . . . .                                                      | 53        |
| 229 | 3.5.2    | Implementation Considerations . . . . .                                             | 58        |
| 230 | 3.5.3    | Final Remarks . . . . .                                                             | 59        |
| 231 | <b>4</b> | <b>Truncated Generalized Cross Entropy for segmentation</b>                         | <b>61</b> |
| 232 | 4.1      | Loss Functions for Multiple Annotators . . . . .                                    | 61        |
| 233 | 4.1.1    | Generalized Cross Entropy . . . . .                                                 | 62        |
| 234 | 4.1.2    | Extension to Multiple Annotators . . . . .                                          | 63        |
| 235 | 4.1.3    | Reliability Maps and Truncated GCE . . . . .                                        | 64        |
| 236 | 4.2      | Chapter overview . . . . .                                                          | 65        |
| 237 | <b>5</b> | <b>Chained deep learning and <math>TGCE_{SS}</math> loss for image segmentation</b> | <b>67</b> |
| 238 | 5.1      | Backbone Architecture . . . . .                                                     | 68        |
| 239 | 5.1.1    | U-Net Architecture . . . . .                                                        | 68        |
| 240 | 5.2      | Reliability Branch . . . . .                                                        | 69        |
| 241 | 5.2.1    | Scalar Reliability Estimation . . . . .                                             | 69        |
| 242 | 5.2.2    | Feature-based Reliability Estimation . . . . .                                      | 70        |
| 243 | 5.2.3    | Pixel-wise Reliability Estimation . . . . .                                         | 71        |
| 244 | 5.3      | Integration with TGCE Loss . . . . .                                                | 71        |
| 245 | 5.3.1    | Regularization terms . . . . .                                                      | 72        |
| 246 | 5.4      | Training Process . . . . .                                                          | 75        |
| 247 | 5.5      | Experimental Setup . . . . .                                                        | 78        |
| 248 | 5.5.1    | Noisy Annotations Generation . . . . .                                              | 78        |
| 249 | 5.5.2    | Missing Labels Emulation . . . . .                                                  | 81        |
| 250 | 5.6      | Results . . . . .                                                                   | 82        |
| 251 | 5.6.1    | Oxford-IIIT Pet Dataset . . . . .                                                   | 85        |
| 252 | 5.6.2    | TNBC Dataset . . . . .                                                              | 87        |

---

|     |                                   |           |
|-----|-----------------------------------|-----------|
| 253 | 5.7 Summary . . . . .             | 90        |
| 254 | <b>6 Final Remarks</b>            | <b>93</b> |
| 255 | 6.1 Conclusions . . . . .         | 93        |
| 256 | 6.2 Future work . . . . .         | 95        |
| 257 | 6.3 Academic Discussion . . . . . | 97        |
| 258 | <b>Bibliography</b>               | <b>99</b> |



## LIST OF FIGURES

|     |            |                                                                        |    |
|-----|------------|------------------------------------------------------------------------|----|
| 260 | <b>1-1</b> | Estimation of the popularity of tasks and medical image types based    |    |
| 261 |            | on recent literature review (count of referenced terms.) . . . . .     | 3  |
| 262 | <b>1-2</b> | AI and Machine Learning in medical imaging recent timeline. . . . .    | 4  |
| 263 | <b>1-3</b> | Example of a histopathological image segmented by multiple             |    |
| 264 |            | annotators, illustrating variations in label assignment. . . . .       | 7  |
| 265 | <b>1-4</b> | Summary diagram for problem Statement . . . . .                        | 11 |
| 266 | <b>1-5</b> | Proposed framework for the approach in [López-Pérez et al., 2024]. .   | 18 |
| 267 | <b>2-1</b> | Overview of digital image concepts and their mathematical              |    |
| 268 |            | representations. The figure shows the main types of digital images     |    |
| 269 |            | (grayscale, color, multispectral, and volumetric), their mathematical  |    |
| 270 |            | representations, and the fundamental processes of sampling and         |    |
| 271 |            | quantization. Example images are included to illustrate each type. .   | 28 |
| 272 | <b>2-2</b> | Histology evolution timeline. (Image from [Mazzarini et al., 2021]). . | 29 |
| 273 | <b>2-3</b> | (Above) Whole slide imaging system by Omnyx for slide                  |    |
| 274 |            | digitization. (Below) Comprehensive digital pathology scan             |    |
| 275 |            | visualization from Omnyx designed to streamline pathologists'          |    |
| 276 |            | diagnostic workflow. (From [Farahani et al., 2015]). . . . .           | 30 |
| 277 | <b>2-4</b> | ROI heatmaps measured experimentally from pathologists view            |    |
| 278 |            | patterns and estimated heatmap from a U-Net model from [Ghezloo        |    |
| 279 |            | et al., 2025]. . . . .                                                 | 32 |

|     |            |                                                                                                             |    |
|-----|------------|-------------------------------------------------------------------------------------------------------------|----|
| 280 | <b>2-5</b> | Common learning paradigms. . . . .                                                                          | 35 |
| 281 | <b>2-6</b> | Comparative Trends of the top two most popular Deep Learning                                                |    |
| 282 |            | Frameworks. . . . .                                                                                         | 37 |
| 283 | <b>2-7</b> | Annotations in the Oxford-IIIT Pet data. From left to right: pet                                            |    |
| 284 |            | image, head bounding box, and trimap segmentation (blue:                                                    |    |
| 285 |            | background region; red: ambiguous region; yellow: foreground                                                |    |
| 286 |            | region). . . . .                                                                                            | 39 |
| 287 | <b>2-8</b> | Different staining techniques obtained from multi-stain breast                                              |    |
| 288 |            | cancer dataset [Weitz et al., 2023]. (a) shows H&E, (b) ER, (c) HER2,                                       |    |
| 289 |            | (d) Ki67 and (e) PGR. (f) shows an example of a Whole Slide Imaging                                         |    |
| 290 |            | (WSI) that was excluded since it contains multiple tissue sections. .                                       | 40 |
| 291 | <b>2-9</b> | Screenshot of the DSA and HistomicsTK web interface while creating                                          |    |
| 292 |            | the crowdsourced annotations for the dataset presented by [Amgad                                            |    |
| 293 |            | et al., 2019]. . . . .                                                                                      | 42 |
| 294 | <b>3-1</b> | Synthetic dataset overview for classification task . . . . .                                                | 54 |
| 295 | <b>3-2</b> | Classification results for the Correlated Chained Gaussian Processes                                        |    |
| 296 |            | for Multiple Annotators (CCGPMA) model using GPflux. . . . .                                                | 55 |
| 297 | <b>3-3</b> | Ground truth and noisy labels emulation for annotators across the                                           |    |
| 298 |            | feature space for regression task. . . . .                                                                  | 55 |
| 299 | <b>3-4</b> | Regression results for the CCGPMA model using GPflux. . . . .                                               | 56 |
| 300 | <b>4-1</b> | Working mechanism of the proposed Truncated Generalized Cross                                               |    |
| 301 |            | Entropy for Semantic Segmentation ( $TGCE_{SS}$ ) loss function. The loss                                   |    |
| 302 |            | combines reliability maps $\Lambda_r$ with the model's predictions $f(\mathbf{X}; \theta)$                  |    |
| 303 |            | to compute a weighted loss that accounts for annotator reliability                                          |    |
| 304 |            | across different image regions. . . . .                                                                     | 65 |
| 305 | <b>5-1</b> | Original U-NET architecture from [Ronneberger et al., 2015]. . . . .                                        | 68 |
| 306 | <b>5-2</b> | Overview of the proposed model architecture. . . . .                                                        | 69 |
| 307 | <b>5-3</b> | Example of the $\mathcal{L}_\lambda(z)$ function for values of $\alpha_\lambda = 0.2$ and $\beta = 1.0$ . . | 74 |

|     |            |                                                                           |    |
|-----|------------|---------------------------------------------------------------------------|----|
| 308 | <b>5-4</b> | Training process of the proposed model overview. The model first          |    |
| 309 |            | undergoes hyperparameter optimization using Bayesian                      |    |
| 310 |            | optimization, followed by end-to-end training with the optimal            |    |
| 311 |            | parameters. Estimated segmentation and reliability maps are               |    |
| 312 |            | computed for each input image, and the loss function is computed          |    |
| 313 |            | for the entire batch. Network parameters are tuned based on the           |    |
| 314 |            | loss function as objective (with minimization target) whereas the         |    |
| 315 |            | hyperparameters $\phi$ tuning are based on the Dice coefficient of the    |    |
| 316 |            | validation data set (with maximization target). . . . .                   | 78 |
| 317 | <b>5-5</b> | Example of a noisy mask generated by naively introducing random           |    |
| 318 |            | noise into a ground truth mask. Morphological consistency is lost. .      | 79 |
| 319 | <b>5-6</b> | Noisy mask generated by enhancing the disturbances in the                 |    |
| 320 |            | encoder layers weights for the Oxford-IIIT Pet Dataset.                   |    |
| 321 |            | Morphological consistency is preserved. From left to right, SNR           |    |
| 322 |            | levels of noise in the encoder layer are 10, 5, 2, 0, -5 dB. Not adding   |    |
| 323 |            | noisy to the encoder layer would reach to an annotator behavior as        |    |
| 324 |            | good as the originally trained base model. . . . .                        | 80 |
| 325 | <b>5-7</b> | Example of missing labels and noisy annotations emulation for the         |    |
| 326 |            | Oxford-IIIT Pet Dataset. The same label rate ( $r_l = 0.7$ ) is preserved |    |
| 327 |            | for all labelers, which means in this case that every labeler is          |    |
| 328 |            | expected to label 70% of the samples, and skip the remaining 30%. .       | 82 |
| 329 | <b>5-8</b> | Segmentation results for baseline model on the Triple-Negative            |    |
| 330 |            | Breast Cancer (TNBC) (Crowd-Seg) dataset after 20 epochs of               |    |
| 331 |            | training. Even though the segmentation algorithm seems to show            |    |
| 332 |            | progress on its way to capture the underlying target structures ot        |    |
| 333 |            | be segmented, since the ground truth annotation is not expressive         |    |
| 334 |            | enough, the Dice coefficient remains below 0.55, which is not             |    |
| 335 |            | considered a good performance. The baseline model remained with           |    |
| 336 |            | overfitting even after combining techniques of: Bayesian                  |    |
| 337 |            | hyperparameter optimization, data augmentation, learning rate             |    |
| 338 |            | scheduling, dataset shuffling, and early stopping. . . . .                | 84 |

|     |             |                                                                                                      |    |
|-----|-------------|------------------------------------------------------------------------------------------------------|----|
| 339 | <b>5-9</b>  | Segmentation results for scalar model on the TNBC (Crowd-Seg)                                        |    |
| 340 |             | dataset after 20 epochs of training. Again, training metrics show a                                  |    |
| 341 |             | constant grow, but not validation metrics, which shows a combined                                    |    |
| 342 |             | issue of overfitting and unreliable ground truth. . . . .                                            | 84 |
| 343 | <b>5-10</b> | Example evaluation of the proposed model in its last trained epoch.                                  |    |
| 344 |             | An accurate segmentation output is delivered as well as scalar                                       |    |
| 345 |             | reliability estimations from the labelers masks. Disturbance                                         |    |
| 346 |             | signal-to-noise ratios are: $[-20, 0, 10]$ dB from left to right. Best                               |    |
| 347 |             | found hyperparameters: learning rate = 0.0011, $q = 0.4$ , noise                                     |    |
| 348 |             | tolerance = 0.2, $\alpha_{reg} = 0.26$ , $\alpha_{entropy} = 0.03$ , $\alpha_{sum} = 0.27$ . . . . . | 86 |
| 349 | <b>5-11</b> | Example evaluation of the proposed model in some given epoch                                         |    |
| 350 |             | (5/50) of training, which demonstrates a quick convergence of the                                    |    |
| 351 |             | model for correctly estimating the reliability of the labelers masks.                                |    |
| 352 |             | An accurate segmentation output is delivered as well as pixel-wise                                   |    |
| 353 |             | reliability estimations from the labelers masks. Disturbance                                         |    |
| 354 |             | signal-to-noise ratios are: $[-20, 0, 10]$ dB from left to right. Best                               |    |
| 355 |             | found hyperparameters: learning rate = 0.0027, $q = 0.9$ , noise                                     |    |
| 356 |             | tolerance = 0.2, $\alpha_{reg} = 0.38$ , $\alpha_{entropy} = 0.02$ , $\alpha_{sum} = 0.42$ . . . . . | 86 |
| 357 | <b>5-12</b> | Label rate distribution across the labelers for the TNBC (Crowd-Seg)                                 |    |
| 358 |             | dataset. . . . .                                                                                     | 88 |
| 359 | <b>5-13</b> | Classes distribution across the labelers for the TNBC (Crowd-Seg)                                    |    |
| 360 |             | dataset. Most common classes are ‘Tumor’ (class 2), ‘Other’ (class 1)                                |    |
| 361 |             | and ‘Stroma’ (class 3). . . . .                                                                      | 88 |
| 362 | <b>5-14</b> | Dice score of the labelers masks against the ground truth for the                                    |    |
| 363 |             | TNBC (Crowd-Seg) dataset. . . . .                                                                    | 89 |
| 364 | <b>5-15</b> | Example of ground truth with odd geometries (second patch from                                       |    |
| 365 |             | top to bottom), which are not always correct or compatible with the                                  |    |
| 366 |             | underlying structures of the patch itself. . . . .                                                   | 90 |
| 367 | <b>5-16</b> | Training curves for the baseline model on the TNBC (Crowd-Seg)                                       |    |
| 368 |             | dataset. . . . .                                                                                     | 90 |

## LIST OF TABLES

|     |            |                                                                                   |    |
|-----|------------|-----------------------------------------------------------------------------------|----|
| 370 | <b>1-1</b> | Summary of state-of-the-art approaches for handling multiple                      |    |
| 371 |            | annotators in medical image segmentation . . . . .                                | 22 |
| 372 | <b>3-1</b> | Comparison of Multiple Annotator Learning Methods . . . . .                       | 57 |
| 373 | <b>5-1</b> | Hyperparameter importance ranking for regularization parameters                   |    |
| 374 |            | obtained from Optuna optimization study after 20 iterations. This                 |    |
| 375 |            | study used preselected values of $q = 0.66$ , $\tilde{C} = 0.62$ , and learning   |    |
| 376 |            | rate $= 6.5 \times 10^{-4}$ . The importance scores indicate the relative impact  |    |
| 377 |            | of each regularization hyperparameter on the model's performance.                 | 76 |
| 378 | <b>5-2</b> | Comprehensive hyperparameter importance ranking including all                     |    |
| 379 |            | parameters obtained from Optuna optimization study after 20                       |    |
| 380 |            | iterations. This study shows the relative importance of mandatory                 |    |
| 381 |            | loss function parameters ( $q$ , $\tilde{C}$ ), learning rate, and regularization |    |
| 382 |            | parameters together. . . . .                                                      | 77 |
| 383 | <b>5-3</b> | Performance metrics (Dice and Jaccard) at different labeling rates.               |    |
| 384 |            | Best values per category are bolded. . . . .                                      | 85 |



## LIST OF ALGORITHMS

|            |                                                                    |    |
|------------|--------------------------------------------------------------------|----|
| <b>5-1</b> | Multi-Annotator Dataset Generation with Noisy and Missing Labels . | 83 |
|------------|--------------------------------------------------------------------|----|





- 388 **CAD** Computer-Aided Diagnosis 2, 5, 8
- 389 **CCGP** Correlated Chained Gaussian Processes 18
- 390 **CCGPMA** Correlated Chained Gaussian Processes for Multiple Annotators xxii, 18, 19,  
391 53, 55–57, 59, 95
- 392 **CE** Cross Entropy 62, 65
- 393 **CGP** Chained Gaussian Processes 18, 45, 46
- 394 **CKA** Centered Kernel Alignment 44
- 395 **CNN** Convolutional Neural Networks 3, 15, 67, 95
- 396 **CR** Crowdsourcing 16, 17
- 397 **CT** Computed Tomography 1, 2, 12
- 398 **ELBO** Evidence Lower Bound 19
- 399 **EM** Expectation-Maximization 15, 20
- 400 **GCE** Generalized Cross Entropy 62, 65, 71
- 401 **GCECDL** Generalized Cross-Entropy-based Chained Deep Learning 21
- 402 **GP** Gaussian Process 46, 58
- 403 **GPMC** Gaussian Process Monte Carlo 48
- 404 **GPR** Gaussian Process Regression 48

- 405 **ISS** Image Semantic Segmentation 2, 5, 8, 10, 14, 21, 23, 37, 93, 95
- 406 **KAAR** Kernel Alignment-Based Annotator Relevance Analysis 43, 44, 46, 57
- 407 **LF** Latent Function 18
- 408 **LKAAR** Localized Kernel Alignment-Based Annotator Relevance Analysis 44–46, 57
- 409 **MAE** Mean Absolute Error 62, 65, 71
- 410 **MITs** Medical Imaging Techniques 1
- 411 **ML** Machine Learning 10, 12, 21
- 412 **MRI** Magnetic Resonance Imaging 1, 2
- 413 **MV** Majority Voting 10, 12, 20, 22
- 414 **OCR** Optical Character Recognition 12
- 415 **PET** Positron Emission Tomography 14
- 416 **RCDNN** Regularized Chained Deep Neural Network 47, 57
- 417 **ROI** Region of Interest 2, 8, 16, 31, 87
- 418 **SGPMC** Sparse Gaussian Process Monte Carlo 48
- 419 **SGPR** Sparse Gaussian Process Regression 48
- 420 **SLFM** Semi-Parametric Latent Factor Model 18, 51, 52
- 421 **SS** Semantic segmentation 3, 67
- 422 **STAPLE** Simultaneous Truth and Performance Level Estimation 13–15, 38
- 423 **TNBC** Triple-Negative Breast Cancer xxiii, xxiv, 41, 84, 87–91
- 424 **WSI** Whole Slide Imaging xxii, 1, 5, 8, 16, 29, 31, 40, 87

425

426

CHAPTER

427

ONE

428

429

INTRODUCTION

430

## 1.1 Motivation

431 Since Roentgen's discovery of X-rays in 1895, medical imaging has advanced  
432 significantly, with modalities like radionuclide imaging, ultrasound, **Computed**  
433 **Tomography (CT)**, **Magnetic Resonance Imaging (MRI)**, and digital radiography  
434 emerging over the past 50 years. Modern imaging extends beyond image  
435 production to include processing, display, storage, transmission and analysis.  
436 [Zhou et al., 2021]. Other **Medical Imaging Techniques (MITs)** have arose during  
437 the last decades, some of them implying only the examination of certain pieces or  
438 tissues instead of complete patients, like histopathological images, which are  
439 images of tissue samples obtained from biopsies or surgical resections and are  
440 widely used for the diagnosis of diseases like cancer through **Whole Slide Imaging**  
441 **(WSI)** scanners [Rashmi et al., 2021].

442 Along with the advances in technologies for medical images acquisition,  
443 computational technologies on pattern recognition and artificial intelligence have

also emerged, allowing the development of **Computer-Aided Diagnosis (CAD)** systems based on Machine Learning algorithms. These systems aim to assist physicians in the diagnosis and treatment of diseases, by providing a second opinion or by automating the analysis of medical images. [Panayides et al., 2020]. One of the most common tasks in which machine learning technologies are being used in the universe of medical images is **Image Semantic Segmentation (ISS)**, which consists of assigning a label to each pixel in an image according to the object it belongs to. This is crucial for the development of **CAD** systems, as it allows the identification of **Region of Interest (ROI)** in the images, which can be used to detect and classify diseases [Azad et al., 2024].

The application of Machine Learning in medical imaging has grown significantly, including classification, segmentation, anomaly detection, super-resolution, image registration, and synthetic image generation [Brito-Pacheco et al., 2025]. Among imaging modalities, X-rays and **CT** scans are widely used for classification and anomaly detection, especially in pulmonary and oncological applications. **MRI** and ultrasound play a crucial role in segmentation and resolution enhancement, while PET/SPECT imaging is essential for anomaly detection in oncology and neurodegenerative diseases [Brito-Pacheco et al., 2025]. Histopathology is rapidly gaining prominence, particularly in segmentation and feature extraction, where AI-driven techniques aid in automated cancer diagnosis and tissue structure analysis. The integration of Deep Learning in histological image processing is revolutionizing pathology, enabling more precise and efficient diagnostics. A brief comparison of the tasks and medical image types based on counted referenced terms occurrences from recent literature review, can be seen in Figure 1-1. [Yu et al., 2025], [Brito-Pacheco et al., 2025], [Ryou et al., 2025], [Hu et al., 2025], [Elhaminia et al., 2025]

Throughout the time that medical images and computational techniques have been developed, a coevolution has occurred, where the complexity of the tasks and the capabilities of the computational techniques have been increasing [Zhou et al., 2021]. Initially, these needs were covered with simple morphological




Figure 1-1 Estimation of the popularity of tasks and medical image types based on recent literature review (count of referenced terms.)

transformations, which implied no training process or elaborated optimization, this is an advantage in terms of development costs but has limited performance. However, as the complexity of the tasks increased, the need for more sophisticated techniques arose, leading to the application of advanced statistical tools and Machine Learning algorithms like Support Vector Machines, Decision Trees, and SGD Neural Networks [Avanzo et al., 2024]. The coevolution of advances in medical image acquisition, computational power (i.e. Moore's law) and statistical/mathematical techniques have led to a convergence for merging state of the art algorithms with medical imaging [Shalf, 2020]. Figure 1-2 shows a timeline of coevolution between some conspicuous advances in computational pattern recognition and its medical applications in different scopes (besides medical imaging) [Avanzo et al., 2024].

Convolutional Neural Networks (CNN) have been widely used in Semantic segmentation (SS), as they have outperformed traditional Machine Learning algorithms both medical and non medical images [Xu et al., 2024] [Sarvamangala and Kulkarni, 2022]. However, most CNN architectures are deep, which introduces the requirement of large amounts of training data. This introduces a




Figure 1-2 AI and Machine Learning in medical imaging recent timeline.

problem since the acquisition and annotation of medical images are expensive and time-consuming processes. This is especially true for ISS tasks, as they require pixel-level annotations, which is taxing in terms of cost, time and the logistics involved [Bhalgat et al., 2018]. Other modalities face this problem through less expensive annotation strategies like bounding boxes or anatomical landmarks for being used in a semi-supervised strategy [Shah et al., 2018].

However, many medical images datasets contain a high variability in class sizes and variations in colors, which is specially noticeable in histopathological images because of the usage of different staining techniques and other factors which can affect the color of the images (see section 2.2.3). This variability can lead to a significant loss of efficiency of Machine Learning models when using a mixed supervision strategy, as the model can be biased towards the most common classes or colors in the dataset [Shah et al., 2018].

This is where other solutions arise to tackle the problem of the weak image annotation while maintaining low costs. One of these solutions is the *crowdsourcing* strategy, which consists of having multiple annotators labeling the same image, and then combining the labels to obtain a consensus label [Lu et al., 2023]. This strategy can lead to a labeling cost reduction when different levels of expertise are combined, since the crowd may be composed of both experts and laymen, the latter being less expensive to hire [López-Pérez et al., 2023].

Recently, diagnosis, prognosis and treatment of cancer have heavily relied on histopathology, where tissue samples are obtained through biopsies or surgical resections and critical information that helps pathologists determine the presence and severity of malignancies [López-Pérez et al., 2024]. The segmentation of histopathological images enables precise identification of structures such as nuclei, glands, and tumors, which are essential for assessing disease progression and treatment response [Rashmi et al., 2021]. Accurate segmentation is particularly crucial in digital pathology, where whole-slide images (WSI) are analyzed using AI-powered CAD systems to support clinical decision-making [López-Pérez et al., 2024].

521 A major challenge in histopathological image segmentation arises from the  
522 variability in annotations provided by different pathologists. Unlike natural  
523 images, where object boundaries are often well-defined, histological structures  
524 may have ambiguous borders, leading to inconsistencies among annotators  
525 [López-Pérez et al., 2023]. Because of this, crowdsourcing labeling is one of the  
526 most popular approaches, as illustrated in Figure 1-3, <sup>1</sup>. These discrepancies  
527 highlight the need for models that can handle annotation uncertainty effectively.  
528 Leveraging crowdsourcing strategies and Machine Learning techniques that infer  
529 annotator reliability can enhance segmentation performance while reducing costs.

---

<sup>1</sup>obtained from a real world Triple Negative Breast Cancer (TNBC) dataset published in [López-Pérez et al., 2023]






Figure 1-3 Example of a histopathological image segmented by multiple annotators, illustrating variations in label assignment.

## 1.2 Problem Statement

Throughout the development of medical technology and CAD, the task of ISS has become a crucial step in delivering precise diagnosis and treatment planning [Giri and Bhatia, 2024]. Particularly, in the area of histopathological studies, the usage of Whole Slide Images (WSI) is rather common since this method delivers high quality imaging and allows for the diagnosis of diseases like cancer [Lin et al., 2024].

ISS task consists of assigning a label to each pixel in an image according to the object it belongs to. Accurate segmentation is essential for the development of CAD systems, as it allows the identification of regions of interest (ROI) in the images, which can be used to detect and classify diseases and hence, treatment planning [Sarvamangala and Kulkarni, 2022]. However, modern computational solutions for ISS tasks involve the use of deep learning, which mostly rely large amounts of labeled data to train the models on supervised learning techniques. This means that the model is trained on a dataset with ground-truth labels, which are assumed to be correct and consistent across all samples. In practice, this assumption is often violated due to the high technical complexity of labeling these segments<sup>2</sup>.

The process of labeling medical images is often managed with the help of specialized software tools that allow the annotators to draw the regions, delivering a standard format for the labeled masks [Habis, 2024]. Despite the help of these tools, the labeling process in WSI can have high costs, as it requires long hours of work from specialized personnel. In the pursuit of reducing costs, the labeling processes is often done by multiple labelers with varying levels of expertise, equalizing the cost of the labeling process. However, this strategy can lead to inconsistent labels, as the consensus between the labelers may not be exact due to the diversity in depth of knowledge and experience of the labelers [Xu et al., 2024]. These inconsistencies are discussed in the subsections 1.2.1 and 1.2.2.

---

<sup>2</sup>compared to a more trivial task like image classification on ordinary an well known classes like MNIST

### 1.2.1 Variability in Expertise Levels

One of the primary sources of inter-observer variability in medical image segmentation is the difference in expertise levels among annotators [López-Pérez et al., 2023]. Experienced radiologists and pathologists tend to produce highly precise annotations, whereas novice labelers may introduce systematic biases due to their limited familiarity with subtle image features. Studies have demonstrated that annotation accuracy *tends* to improve with experience, yet medical institutions often rely on a mix of annotators to manage costs and workload distribution [Lu et al., 2023].

The training background of annotators and institutional guidelines play a crucial role in shaping labeling practices. Different medical schools and hospitals may adopt distinct segmentation protocols, leading to inconsistencies when datasets are combined from multiple sources [López-Pérez et al., 2023]. For example, some institutions may emphasize conservative delineation of tumor boundaries, while others adopt a more inclusive approach. Such variations contribute to systematic biases in medical image datasets [Banerjee et al., 2025].

Medical images frequently contain complex structures which are difficult to distinguish, making segmentation inherently subjective. For instance, tumor margins in histopathological slides may not have well-defined edges, leading to variations in how different annotators delineate the regions of interest [Carmo et al., 2025]. These discrepancies arise not only from technical expertise but also from differences in perception and interpretation.

### 1.2.2 Technical Constraints and Image Quality

Technical constraints in medical imaging, such as resolution differences, noise levels, and contrast variations, can significantly impact segmentation accuracy.

Lower-resolution images may obscure fine structures, leading to inconsistencies in boundary delineation [Zhou et al., 2024].

When combined with long sessions, images with insufficient quality might also increase the cognitive load of the annotators, leading to fatigue and reduced precision in labeling [Kim et al., 2024]. This is particularly relevant in histopathological studies, where the staining process and tissue preparation can introduce color variations and artifacts that affect image quality, even if the same scanning equipment is used [Karthikeyan et al., 2023].

### 1.2.3 Research Question

Given the challenges posed by inconsistent labels in medical image segmentation, this work aims to address the following research question:

#### Research Question

*How can we develop a learning approach for ISS tasks in medical images that can adapt to inconsistent labels without requiring explicit supervision of labeler performance, while addressing challenges related to variability in expertise levels and technical constraints, and maintaining interpretability and generalization?*

## 1.3 State of the Art

Certainly, in general Machine Learning (ML) classification tasks<sup>3</sup> where multiple annotators are involved, Majority Voting (MV) is by far the simplest possible approach to implement. This concept was born multiple times and divergently in

---

<sup>3</sup>In this work, image segmentation is considered as a particular case of classification in which target classes are assigned pixel-wise.




Figure 1-4 Summary diagram for problem Statement

multiple fields, but it was described as relevant for ML and pattern recognition labeling for classification in [Lam and Suen, 1997], in which the approach is exposed as simple, yet powerful. The authors describe the MV as a method that can be used to improve the accuracy of classification tasks by combining the labels of multiple annotators. The method is based on the assumption that the majority vote of the annotators is more likely to be correct than the vote of a single annotator. The authors also describe the method as a straightforward way to improve the accuracy of classification tasks without the need for complex algorithms or additional data. The authors also prove this method to deliver very similar results to more complicated approaches (Bayesian, logistic regression, fuzzy integral, and neural network) in the particular task of Optical Character Recognition (OCR). Despite its simplicity, modern solutions for delivering accurate medical image segmentation models still rely on Majority Voting at some stage. In [Elnakib et al., 2020], a majority voting strategy is used for delivering a final output based on the labels of multiple models (VGG16-Segnet, Resnet-18 and Alexnet) in CT images for Liver Tumor Segmentation. Majority voting as a technique for label aggregation is a powerful approach for its simplicity in many use cases in which the target to be labeled is not tied to an expertise related task, otherwise, the assumption of equal expertise among the labelers may introduce bias in the final label, which is not desirable in the case of highly technical annotations like medical images. In subsection 1.3.1, a review is done on literature which no longer assumes the naive approach of equal expertise among labelers and face the challenge of learning from inconsistent labels will be presented.

### 1.3.1 Facing Annotation Variability in Medical images

Learning from crowds approaches in general face the challenge of not having a ground truth label and hence, an intrinsic difficulty in measuring the real reliability of the labelers annotations. Some approaches assume beforehand a certain level of expertise for each labeler based on experience as an input, like in [TIAN and Zhu,

2015], which introduce the concept of max margin majority voting, using the reliability vector as weights for the binary and multiclass classifier. The crowdsourcing margin is the minimal difference between the aggregated score of the potential true label and the scores for other alternative labels. Accordingly, the annotators' reliability is estimated as generating the largest margin between the potential true labels and other alternatives. The problem introduced in this approach is assuming an stationary reliability per expert across the whole input space, which is imprecise since annotators performance may change between different tasks or even between different regions of the same image.

### STAPLE Mechanism

The Simultaneous Truth and Performance Level Estimation (STAPLE) algorithm, introduced in [Warfield et al., 2004] is a probabilistic framework that estimates a hidden true segmentation from multiple segmentations provided by different raters. It also estimates the reliability of each rater by computing their sensitivity and specificity.

The STAPLE algorithm's goal is to maximize the log likelihood function:

$$(\hat{\mathbf{p}}, \hat{\mathbf{q}}) = \arg \max_{\mathbf{p}, \mathbf{q}} \ln f(\mathbf{D}, \mathbf{T} \mid \mathbf{p}, \mathbf{q}). \quad (1-1)$$

Where  $\mathbf{D}$  is the set of segmentations provided by the raters,  $\mathbf{T}$  is the hidden true segmentation,  $p$  is the sensitivity and  $q$  is the specificity of the raters.

This is achieved by using the Expectation-Maximization algorithm to maximize the log likelihood function, which is done iteratively with step computations:

$$\begin{aligned}
(p_j^{(k)}, q_j^{(k)}) = \arg \max_{p_j, q_j} & \sum_{i: D_{ij}=1} W_i^{(k-1)} \ln p_j \\
& + \sum_{i: D_{ij}=1} (1 - W_i^{(k-1)}) \ln(1 - q_j) \\
& + \sum_{i: D_{ij}=0} W_i^{(k-1)} \ln(1 - p_j) \\
& + \sum_{i: D_{ij}=0} (1 - W_i^{(k-1)}) \ln q_j.
\end{aligned} \tag{1-2}$$

646 The weight terms  $W_i^{(k-1)}$  and  $(1 - W_i^{(k-1)})$  represent the current estimates of the  
647 probability that voxel  $i$  belongs or does not belong to the true segmentation from  
648 the previous E-step. The log-likelihood contributions  $\ln p_j$ ,  $\ln(1 - q_j)$ ,  $\ln(1 - p_j)$ , and  
649  $\ln q_j$  account for true positives, false positives, false negatives, and true negatives,  
650 respectively, weighted by the current estimate of true segmentation probabilities.

651 The capacity of STAPLE to accurately estimate the true segmentation, even in the  
652 presence of a majority of raters generating correlated errors, was demonstrated,  
653 which makes it theoretically a strong choice for setting a ground-truth in binary or  
654 multiclass medical ISS tasks.

655 The popularity and performance of STAPLE has led to its usage in modern  
656 applications medical image, 3d spatial images due to its assumption of decision  
657 space being based on voxel-wise decisions, like the authors in [Grefve et al., 2024]  
658 which applied the algorithm on Positron Emission Tomography (PET) images.  
659 Other authors still rely heavily on STAPLE for setting a ground truth consensus for  
660 histopathological images, like [Qiu et al., 2022].

661 However, the STAPLE algorithm has some limitations. It assumes independent  
662 rater errors, which may not hold in practice, leading to biased estimates. STAPLE  
663 is also sensitive to low-quality annotations, potentially degrading final  
664 segmentations if the weights are not initialized correctly. The algorithm tends to  
665 over-smooth results, blurring fine details, and struggles with multi-class



segmentation. Computationally, it is expensive due to its iterative Expectation-Maximization (EM) approach. Additionally, STAPLE cannot correct systematic biases in annotations and depends on initial estimates, impacting accuracy. Lastly, the estimated performance levels lack interpretability, making it difficult to assess annotator reliability effectively.

Finally, while STAPLE proves valuable for aggregating existing labels and supporting other methodologies, its primary limitation lies in its inability to generate annotations for novel, unlabeled images. This restricts its applicability to scenarios where multiple annotations already exist, rather than serving as a standalone solution for new image segmentation tasks.

## U-shaped CNNs

Since the introduction of U-Net [Ronneberger et al., 2015] in 2015 for biomedical image segmentation, U-shaped CNNs have become a prevalent architecture in medical image segmentation tasks. The U-Net's success stems from its ability to capture both global and local information through its contracting and expanding paths, making it particularly effective for complex and heterogeneous structures, even with limited annotated data. This architecture has been successfully applied to various medical image segmentation tasks, including organ segmentation, tumor segmentation, and brain structure segmentation [Azad et al., 2024].

The U-Net architecture consists of a symmetric encoder-decoder structure with skip connections. The encoder path progressively reduces spatial dimensions while increasing feature channels through a series of convolutional and max-pooling layers, capturing high-level semantic information. The decoder path uses transposed convolutions to gradually recover spatial resolution while reducing feature channels. Skip connections between corresponding encoder and decoder layers preserve fine-grained details by concatenating high-resolution features from the encoder with upsampled features in the decoder, enabling precise localization of structures.

In [López-Pérez et al., 2024] two networks are trained for delivering a final segmentation. One network is trained to estimate the annotators reliability and another one is trained to segment the image. The first network is a Deep Neural Network that takes as input features of image and the one-hot encoded identifiers of the labelers and outputs a reliability map across the image feature space. This map is then used to weight the contribution of each annotator to the final segmentation. The second network is the U-Net used for segmentation.

In this approach, it is assumed that the images are labeled for at least one labeler and not all of them, which is closer to a real world scenario, in which images with variability in the amount of annotations per patch is common. Hence, the input data can be modeled as:

$$\mathcal{D} = (\mathbf{X}, \tilde{\mathbf{Y}}) = \{(\mathbf{x}_n, \tilde{\mathbf{y}}_n^r) : n = 1, \dots, N; r \in R_n\}, \quad (1-3)$$

Where every  $\mathbf{x}_n$  is an input patch from a ROI in one WSI,  $\tilde{\mathbf{y}}_n$  is the noisy annotation from the  $r$  labeler,  $N$  is the number of patches in the dataset and  $R_n \subset \{1, \dots, R\}$  is the set of labelers that annotated the image  $\mathbf{x}_n$ .

The authors then assume the annotator network to deliver a reliability map  $\{\hat{\mathbf{A}}_\phi^{(r)}(\mathbf{x})\}_{r \in R_n}$  with different dimensions:

- **Crowdsourcing (CR) global:** a single reliability vector per labeler with dimensions  $C$  which represent global reliability of the labeler across all input space.
- **CR image:** a single reliability vector per image per labeler with dimensions  $C$  which represent local reliability of the labeler across the image.
- **CR pixel:** a reliability matrix per image per labeler, with dimensions  $C$  which represent local reliability of the labeler across all the pixels in the image.

717 These differences in dimensions are determined by the feature extraction space  
 718 from segmentation network which feed the input of the annotator network, which  
 719 the authors vary for experimentation purposes.

720 Let  $\mathbf{p}_\theta(\mathbf{x}_n)$  be the estimated latent (ground truth) segmentation produced by the  
 721 segmentation U-Net network. The estimated segmentation probability mask for  
 722 each annotator is then computed as the product:

$$\mathbf{p}_{\theta,\phi}^{(r)}(\mathbf{x}_n) := \mathbf{A}_\phi^{(r)}(\mathbf{x}) \odot \mathbf{p}(\mathbf{x}_n), \quad (1-4)$$

723 where  $\odot$  is the element-wise product and  $\phi$  and  $\theta$  are the parameters of the  
 724 annotator network and the segmentation UNet network, respectively, being the  
 725 latter initialized with a ResNet34 backbone pre-trained on ImageNet.

726 The authors propose a loss function involving cross-entropy and a trace based  
 727 regularization on the reliability map, originally proposed in [Zhang et al., 2020]  
 728 which combined, looks like:

$$\mathcal{L}(\theta, \phi) := \sum_{n=1}^N \sum_{r=1}^R \mathbb{I}(\tilde{\mathbf{y}}_n^{(r)} \in R_n) \cdot \left[ \text{CE} \left( \mathbf{A}_\phi^{(r)}(\mathbf{x}_n) \cdot \mathbf{p}_\theta(\mathbf{x}_n), \tilde{\mathbf{y}}_n^{(r)} \right) + \lambda \cdot \text{tr} \left( \mathbf{A}_\phi^{(r)}(\mathbf{x}_n) \right) \right] \quad (1-5)$$

729 Being  $\mathbb{I}$  the indicator function, CE the cross-entropy loss, and  $\lambda$  the regularization  
 730 parameter.

731 When evaluated on a Triple Negative Breast Cancer dataset, this approach achieves  
 732 a Dice coefficient of 0.7827, outperforming STAPLE (0.7039) and matching  
 733 expert-supervised performance (0.7723). The CR image reliability modeling  
 734 proved most effective, on the other hand CR pixel, while potentially offering  
 735 finer-grained reliability estimation, requires significantly more training data.

736 Despite the proven performance of the approach by the authors, solving the problem  
 737 of multiple labelers with two networks can be overwhelming for the optimization  
 738 process, requiring large amounts of annotated data to properly codify the annotators  
 739 spatial reliabilities, which could be managed by a single model with an appropriate  
 740 loss function.




Figure 1-5 Proposed framework for the approach in [López-Pérez et al., 2024].

## Bayesian Models

Bayesian approaches are a good choice for handling label noise and uncertainty in the labelers. In [Gil-González and Álvarez Meza, 2023] the authors propose a novel approach based on Gaussian Processes for modeling the relationship between the annotators' reliability and the input data, while also preserving the interdependencies among the annotators. This is achieved by introducing Correlated Chained Gaussian Processes for Multiple Annotators (CCGPMA), a framework based on the well known Chained Gaussian Processes (CGP). CGP on itself cannot consider inter-annotator dependencies, thus, the authors introduce the Correlated Chained Gaussian Processes (CCGP) thus, the authors introduce the CCGP to model correlations between the GP latent functions, which are supposed to be generated from a Semi-Parametric Latent Factor Model (SLFM):

$$f_j(\mathbf{x}_n) = \sum_{q=1}^Q w_{j,q} \mu_q(\mathbf{x}_n), \quad (1-6)$$

where  $f_j : \mathcal{X} \rightarrow \mathbb{R}$  is a Latent Function (LF),  $\mu_q(\cdot) \sim \mathcal{GP}(0, k_q(\cdot, \cdot))$  with  $k_q : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  being a kernel function, and  $w_{j,q} \in \mathbb{R}$  is a combination coefficient ( $Q \in \mathbb{N}$ ). This leads to a joint distribution of the form:

$$p(\mathbf{y}, \hat{\mathbf{f}}, u | \mathbf{X}) = p(\mathbf{y} | \boldsymbol{\theta}) \prod_{j=1}^J p(\mathbf{f}_j | \mathbf{u}) p(\mathbf{u}), \quad (1-7)$$

756 where  $\mathbf{y}$  is the vector of noisy labels,  $\hat{\mathbf{f}}$  is the vector of latent functions,  $u$  represents  
757 the inducing points, and  $\mathbf{X}$  is the input data.

758 Combined with inducing-variables based methods for sparse GP approximations,  
759 and maximizing an **Evidence Lower Bound (ELBO)** for the estimation of the  
760 variational parameters, the authors reach a model whose variational expectations  
761 are not analytically tractable, and hence, the authors derive a Gaussian-Hermite  
762 quadrature approach.

763 Finally, the authors extend this approach for being applied to classification and  
764 regression, reaching the only known approach to involve chained gaussian  
765 processes in multiple annotators classification and regression tasks while  
766 preserving the interdependencies among the annotators, and also outperforming  
767 GPC-MV<sup>4</sup>, MA-LFC-C<sup>5</sup>, MA-DGRL<sup>6</sup>, MA-GPC<sup>7</sup>, MA-GPCV<sup>8</sup>, MA-DL<sup>9</sup>, KAAR<sup>10</sup>.

768 **CCGPMA** on itself proposes a good approach for handling label noise and  
769 uncertainty in the labelers for regression and classification tasks, while also  
770 preserving the interdependencies among the annotators, however, it does not face  
771 the image segmentation problem, which is the main focus of this work. Besides,  
772 handling so many latent functions during the optimization process is  
773 computationally expensive, making it on itself infeasible for large and high  
774 resolution datasets.

---

<sup>4</sup>A GPC using the MV of the labels as the ground truth.

<sup>5</sup>A LRC with constant parameters across the input space.

<sup>6</sup>A multi-labeler approach that considers as latent variables the annotator performance.

<sup>7</sup>A multi-labeler GPC, which is an extension of MA-LFC.

<sup>8</sup>An extension of MA-GPC that includes variational inference and priors over the labelers' parameters.

<sup>9</sup>A Crowd Layer for DL, where the annotators' parameters are constant across the input space.

<sup>10</sup>A kernel-based approach that employs a convex combination of classifiers and codes labelers dependencies.

### 1.3.2 Facing Noisy Annotations and Low-Quality Data

The problem of low-quality data and noisy annotations has been tackled with various strategies. One such approach is the use of deep learning models that incorporate loss functions designed to mitigate the effects of unreliable labels. Traditional methods such as *MV* or *EM* have been widely used for aggregating multiple annotators' inputs. However, they assume a homogeneous reliability of annotators, which may not hold in real-world scenarios.

#### Loss Functions in Deep Learning Models

Loss functions are fundamental components in deep learning models that quantify how well a model's predictions match the ground truth. They serve as the objective function that guides the learning process by measuring the discrepancy between predicted and actual values. In classification tasks, the most common loss functions are Cross-Entropy (CE) and Mean Absolute Error (MAE) [Zhang and Sabuncu, 2018]. CE is particularly effective for classification as it heavily penalizes confident but wrong predictions, though it can be sensitive to noisy labels. MAE, on the other hand, is more robust to outliers and assigns equal weights to all mistakes, but typically requires more training iterations. For image segmentation tasks, specialized loss functions have been developed to handle the unique challenges of pixel-wise classification. The Dice loss, which measures the overlap between predicted and ground truth regions, is widely used in medical image segmentation [Zhao et al., 2020]. More recently, the Generalized Cross Entropy (GCE) loss has emerged as a robust alternative that combines the benefits of both CE and MAE, allowing for better handling of noisy labels through a tunable parameter that controls sensitivity to outliers. However, to the best of our knowledge, none of the current loss functions approaches have focused on the multiple labelers scenario, neither on the per labeler spatial reliability estimation.

## Generalized Cross-Entropy for Multiple Annotators Classification

A more recent approach was proposed by [Triana-Martinez et al., 2023], introducing a Generalized Cross-Entropy-based Chained Deep Learning (GCECDL) framework. This method addresses the limitations of traditional label aggregation techniques by modeling each annotator’s reliability as a function of the input data. The approach effectively mitigates the impact of noisy labels by using a noise-robust loss function, balancing Mean Absolute Error (MAE) and Categorical Cross-Entropy (CE). Unlike prior approaches, GCECDL accounts for the dependencies among annotators while encoding their non-stationary behavior across different data samples. Their experiments on multiple datasets demonstrated superior predictive performance compared to state-of-the-art methods, particularly in cases where annotations were highly inconsistent.

The strategy of the authors effectively unlocks the potential of ML models to handle low-quality data and noisy annotations, but it is bounded to classifications tasks only, not being by itself applicable to segmentation tasks. The TGCE equation for handling multiple annotators is defined as:

$$\text{TGCE}(\mathbf{y}, f(\mathbf{x}); \tilde{\lambda}_x, \tilde{C}) = \tilde{\lambda}_x \frac{1 - (\mathbf{1}^\top (\mathbf{y} \odot f(\mathbf{x})))^q}{q} + (1 - \tilde{\lambda}_x) \frac{1 - (\tilde{C})^q}{q}, \quad (1-8)$$

where  $\tilde{\lambda}_x$  represents the annotator reliability,  $\tilde{C}$  is a constant,  $q$  is a parameter that controls the balance between MAE and CE behavior,  $\mathbf{y}$  is the annotation vector, and  $f(\mathbf{x})$  is the model prediction. This approach is more deeply discussed in chapter 4.

## 1.4 Aims

### 1.4.1 General Aim

The main purpose of this work is to adaptively infer the best possible segmentation in ISS tasks for medical images through a novel approach, without needing separate

**Table 1-1** Summary of state-of-the-art approaches for handling multiple annotators in medical image segmentation

| Approach                                                  | Key Characteristics                                                        | Advantages                                                                                                           | Limitations                                                                                                                        |
|-----------------------------------------------------------|----------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| MV [Lam and Suen, 1997]                                   | Simple aggregation of multiple annotators' labels                          | - Simple to implement - No complex algorithms required - Proven effective in many cases                              | - Assumes equal expertise<br>- Sensitive to correlated errors<br>- May introduce bias in expert tasks                              |
| STAPLE [Warfield et al., 2004]                            | Probabilistic framework estimating true segmentation and rater reliability | - Handles binary and multiclass segmentation - Estimates rater sensitivity/specificity - Robust to correlated errors | - Computationally expensive<br>- Sensitive to low-quality annotations<br>- Tends to over-smooth results - Limited interpretability |
| U-Net based approaches [Ronneberger et al., 2015]         | Deep Learning architecture with encoder-decoder structure                  | - Captures both global and local information - Effective for complex structures - Works well with limited data       | - Requires large amounts of training data<br>- Complex optimization process - May need specialized loss functions                  |
| Bayesian Models (CCGPMA) [Gil-Gonzalez et al., 2023]      | Gaussian Processes modeling annotator reliability                          | - Preserves annotator interdependencies - Handles label noise effectively - Works for classification and regression  | - Computationally expensive<br>- Not directly applicable to segmentation<br>- Complex optimization process                         |
| Generalized Cross-Entropy (GCE) [Zhang and Sabuncu, 2018] | Loss function balancing MAE and CE                                         | - Robust to noisy labels - Accounts for non-stationary behavior - Handles annotator dependencies                     | - Limited to classification tasks<br>- Requires parameter tuning<br>- May need adaptation for segmentation                         |



824 inputs about the labeler while offering an interpretable reliability space per labeler  
825 and properly facing the performance inconsistency across the annotators space and  
826 the variability of images quality.

## 827 1.4.2 Specific Aims

- 828 • To develop a novel loss function for ISS tasks in medical images, which  
829 accurately infers the underlying ground truth segmentation mask without  
830 needing separate inputs about the labeler performance.
- 831 • To introduce a tensor map which codifies the reliability of each labeler based  
832 only on the inputs, allowing the model to implicitly weigh the labelers based  
833 on their performance across the mask and classes space.
- 834 • To develop and test a deep learning model for ISS tasks in medical images,  
835 which can learn from inconsistent labels and improve the segmentation  
836 performance compared to other solutions in state of the art.



837

838

CHAPTER

839

TWO

840

841

## CONCEPTUAL PRELIMINARIES

842

### 2.1 Modern Concept of Digital Image

843

844

845

846

847

A digital image is a numerical representation of a visual scene, captured through various imaging devices and stored in a computer. From a mathematical perspective, a digital image can be represented as a function  $f(x, y)$  that maps spatial coordinates  $(x, y)$  to intensity values. In the discrete domain, this function is sampled at regular intervals, creating a matrix of values known as pixels (picture elements).

848

#### 2.1.1 Types of Digital Images

849

##### Grayscale Images

850

851

852

853

854

Grayscale images are the simplest form of digital images, where each pixel represents a single intensity value. Mathematically, a grayscale image can be represented as a 2D matrix  $I$  of size  $M \times N$ , where each element  $I(i, j)$  represents the intensity at position  $(i, j)$ . The intensity values typically range from 0 (black) to 255 (white) in 8-bit images, or from 0 to 65535 in 16-bit images.

## Color images

Color images extend the grayscale concept by representing each pixel with multiple channels, typically Red, Green, and Blue (RGB). A color image can be represented as a 3D matrix  $I$  of size  $M \times N \times 3$ , where  $I(i, j, k)$  represents the intensity of the  $k$ -th color channel at position  $(i, j)$ . Other color spaces like HSV (Hue, Saturation, Value) or CMYK (Cyan, Magenta, Yellow, Key) are also commonly used in different applications.

## Multispectral images

Multispectral images capture information across multiple wavelength bands beyond the visible spectrum. These images can be represented as a 3D matrix  $I$  of size  $M \times N \times B$ , where  $B$  is the number of spectral bands. Each band  $I(i, j, b)$  represents the intensity at position  $(i, j)$  for the  $b$ -th spectral band. This representation is particularly useful in medical imaging, remote sensing, and several scientific applications.

## 3D Images and Volumetric Data

Three-dimensional images extend the concept of pixels to voxels (volume elements). A 3D image can be represented as a 3D matrix  $V$  of size  $M \times N \times D$ , where  $D$  represents the depth dimension. Each voxel  $V(i, j, k)$  represents the intensity at position  $(i, j, k)$  in the 3D space. This representation is fundamental in medical imaging (CT, MRI), scientific visualization, and computer graphics.

### 2.1.2 Mathematical Representations

The mathematical foundation of digital images relies on several key concepts:

- 877 • **Sampling:** The process of converting a continuous image into a discrete  
878 representation. According to the Nyquist-Shannon sampling theorem, the  
879 sampling frequency must be at least twice the highest frequency present in  
880 the image to avoid aliasing.
- 881 • **Quantization:** The process of converting continuous intensity values into  
882 discrete levels. The number of quantization levels determines the image's  
883 bit depth and affects its quality and storage requirements.
- 884 • **Resolution:** The number of pixels per unit length in an image, typically  
885 measured in pixels per inch (PPI) or dots per inch (DPI).
- 886 • **Dynamic range:** The ratio between the maximum and minimum measurable  
887 light intensities in an image, often expressed in decibels (dB).

888 The mathematical representation of a digital image can be expressed as:

$$I(x, y) = \sum_{i=0}^{M-1} \sum_{j=0}^{N-1} f(i, j) \cdot \delta(x - i, y - j) \quad (2-1)$$

889 where  $I(x, y)$  is the digital image,  $f(i, j)$  represents the intensity values, and  $\delta(x -$   
890  $i, y - j)$  is the Kronecker delta function [Duarte and Baraniuk, 2012].

891 For color images, the representation extends to:

$$I(x, y) = \begin{bmatrix} I_R(x, y) \\ I_G(x, y) \\ I_B(x, y) \end{bmatrix} \quad (2-2)$$

892 where  $I_R$ ,  $I_G$ , and  $I_B$  represent the red, green, and blue channels respectively.




Figure 2-1 Overview of digital image concepts and their mathematical representations. The figure shows the main types of digital images (grayscale, color, multispectral, and volumetric), their mathematical representations, and the fundamental processes of sampling and quantization. Example images are included to illustrate each type.

## 2.2 Digital Histopathological Images

Digital histopathology represents a significant advancement in medical imaging, where traditional glass slides containing tissue samples are digitized using specialized scanning devices. This transformation has revolutionized the field of pathology by enabling remote diagnosis, computer-aided analysis, and digital archiving of tissue samples [Amgad et al., 2019]. This process has evolved significantly over the past few decades, as shown in Figure 2-2.




Figure 2-2 Histology evolution timeline. (Image from [Mazzarini et al., 2021]).

### 2.2.1 Whole Slide Imaging (WSI)

Whole Slide Imaging (WSI) is the process of digitizing entire glass slides at high resolution, creating a digital representation that can be viewed, analyzed, and shared electronically. Modern WSI scanners use sophisticated optical systems that capture multiple fields of view at high magnification, which are then stitched

905 together to create a seamless digital image [Hu et al., 2025]. These systems  
906 incorporate high-resolution objectives with magnifications ranging from 20x to  
907 40x, precise motorized stages for accurate slide positioning, automated focus  
908 systems to maintain image quality, and high-quality cameras equipped with large  
909 sensor arrays. The resulting digital slides can reach sizes of several gigabytes,  
910 containing billions of pixels that capture the microscopic details of tissue samples  
911 [Hu et al., 2025]. Figure 2-3 shows a whole slide imaging system by Omnyx for  
912 slide digitization and a comprehensive digital pathology interface from Omnyx  
913 designed to streamline pathologists' diagnostic workflow [Farahani et al., 2015].



Figure 2-3 (Above) Whole slide imaging system by Omnyx for slide digitization. (Below) Comprehensive digital pathology scan visualization from Omnyx designed to streamline pathologists' diagnostic workflow. (From [Farahani et al., 2015]).



### 2.2.2 Regions of Interest (ROI)

In digital histopathology, **Region of Interests (ROIs)** are specific areas within a whole slide image that contain diagnostically relevant information. The importance of **ROIs** lies in their ability to focus computational analysis on relevant areas, reduce computational complexity in automated systems, facilitate targeted diagnosis and research, and enable efficient storage and transmission of critical information.

Selecting **ROIs** is a fundamental step in medical image analysis across all imaging modalities, where an ROI represents a subset of an image appropriate for the intended analysis and is typically identified manually by experts [Hossain et al., 2023]. In modern pathology, the analysis involves processing multidimensional and high-resolution **WSI** tiles automatically, containing an overwhelming quantity of structural and functional information. Despite recent improvements in computing capacity, analyzing such extensive data remains challenging but vital for accurate analysis. Automatic ROI detection can significantly reduce the number of pixels to be processed, speed up the analysis, improve accuracy, and reduce dependency on pathologists [Hossain et al., 2023].

The challenge of acquiring image domain annotations for ROI detection is particularly significant due to the knowledge and commitment required of expert pathologists. Pathologists often identify regions in whole slide images with diagnostic relevance rather than examining the entire slide, with a positive correlation between the time spent on these critical image regions and diagnostic accuracy [Ghezloo et al., 2025]. This selective viewing behavior highlights the importance of developing automated systems that can replicate pathologists' diagnostic patterns.

Recent advances in deep learning techniques have offered improvements in computer-aided diagnosis systems for ROI detection. Innovative approaches include generating heatmaps to represent pathologists' viewing patterns during

diagnosis and using these to guide deep learning architectures during training [Ghezloo et al., 2025]. Such systems can integrate pathologists' domain expertise to enhance region of interest detection without needing individual case annotations, potentially surpassing traditional approaches based on color and texture image characteristics.

The clinical application of automated ROI detection has been demonstrated in specific diagnostic contexts, such as Human Epidermal Growth Factor Receptor 2 (HER2) grading for breast cancer patients. Existing HER2 grading relies on manual ROI selection, which is tedious, time-consuming, and suffers from inter-observer and intra-observer variability [Hossain et al., 2023]. Studies have found that HER2 grade changes with ROI selection, emphasizing the critical importance of accurate and consistent ROI detection methods. Advanced approaches using Vision Transformer architectures have shown promising results, achieving high accuracy rates for ROI detection while significantly reducing processing time and improving diagnostic agreement with clinical scores [Hossain et al., 2023].

These automated ROI detection methods not only enhance diagnostic accuracy and efficiency but also streamline annotation processes and aid in the training of new pathologists. By incorporating pathologists' expertise through heatmap generation and advanced deep learning architectures, these systems demonstrate the potential to revolutionize digital pathology workflows while maintaining the quality and reliability of diagnostic assessments [Ghezloo et al., 2025] (Figure 2-4).




Figure 2-4 ROI heatmaps measured experimentally from pathologists view patterns and estimated heatmap from a U-Net model from [Ghezloo et al., 2025].

### 2.2.3 Staining Techniques

Histopathological analysis relies heavily on various staining techniques to enhance the visibility of different tissue components and cellular structures. The choice of staining method depends on the specific diagnostic requirements and the type of tissue being examined.

#### Hematoxylin and Eosin (H&E)

Hematoxylin and Eosin (H&E) staining is the most widely used technique in histopathology, particularly in breast cancer diagnosis [Pan et al., 2021]. This staining method provides essential visualization through two components: hematoxylin, which stains cell nuclei blue/purple to highlight nuclear morphology, and eosin, which stains cytoplasm and extracellular matrix pink/red to reveal tissue architecture.

The popularity of H&E staining in breast cancer histopathology stems from its ability to clearly visualize tumor architecture and growth patterns, distinguish between different types of breast cancer, identify important diagnostic features like nuclear pleomorphism, and assess tumor grade and stage. Beyond breast cancer, H&E staining finds extensive application across various medical specialties including general pathology, dermatology, gastroenterology, neurology, and oncology.

#### Special Stains

In addition to H&E, various special stains are used for specific diagnostic purposes. Immunohistochemistry (IHC) uses antibodies to detect specific proteins, playing a crucial role in subtyping breast cancers. Key IHC stains include Estrogen Receptor (ER) staining for detecting estrogen receptors, Progesterone Receptor (PGR) staining for assessing progesterone receptor status, Human Epidermal Growth

Factor Receptor 2 (HER2) staining for evaluating HER2 protein expression, and Ki67 staining for measuring cellular proliferation rates. These markers are particularly crucial in breast cancer diagnosis and treatment planning, as they help determine the molecular subtype of the cancer and guide personalized therapeutic approaches. Other specialized stains include Periodic Acid-Schiff (PAS) for highlighting carbohydrates and basement membranes, Masson's Trichrome for distinguishing between collagen and muscle fibers, and silver stains for detecting microorganisms and nerve fibers. These specialized staining techniques complement H&E by providing additional diagnostic information that is crucial for accurate diagnosis and treatment planning [Weitz et al., 2023]. Examples of these staining techniques are shown in Figure 2-8.

## 2.3 Deep Learning Fundamentals

Deep Learning has emerged as a powerful subset of Machine Learning, revolutionizing the field of artificial intelligence. Its roots can be traced back to the early development of artificial neural networks in the 1940s and 1950s, with significant milestones including the perceptron in 1958 and the backpropagation algorithm in the 1980s. However, it wasn't until the early 21st century, with the advent of more powerful computational resources and the availability of large datasets, that deep learning truly began to flourish [Werbos, 1990].

### 2.3.1 Learning Paradigms

Machine Learning systems can be categorized into three main learning paradigms. The most common approach is supervised learning, where models learn from labeled data by mapping inputs to known outputs. This paradigm requires a large amount of labeled training data, which can be expensive and time-consuming to

1012 acquire. Semi-supervised learning offers a hybrid approach that leverages both  
 1013 labeled and unlabeled data [Alloghani et al., 2020], proving particularly useful  
 1014 when labeled data is scarce but unlabeled data is abundant. Finally, unsupervised  
 1015 learning enables models to discover patterns and structures from unlabeled data  
 1016 without explicit guidance [Alloghani et al., 2020], making it valuable for tasks like  
 1017 clustering and dimensionality reduction. Being deep learning a subset of Machine  
 1018 Learning, it is based on the same paradigms. An illustration of the three main  
 1019 learning paradigms is shown in Figure 2-5.




Figure 2-5 Common learning paradigms.

## 1020 2.3.2 Architecture and Training

1021 Deep Learning architectures are characterized by their layered structure, where  
 1022 each layer progressively extracts and transforms features from the input data  
 1023 [Taye, 2023]. The early layers typically focus on low-level feature extraction, such  
 1024 as edges, textures, and basic patterns in the case of image processing [Taye, 2023].  
 1025 As information flows through the network, middle layers combine these basic  
 1026 features into more complex representations. The final layers perform high-level  
 1027 reasoning and make the ultimate predictions or classifications.

1028 The training process commonly employs gradient descent as its primary  
1029 optimization strategy, which iteratively adjusts the model's parameters to  
1030 minimize a loss function. However, other optimization algorithms like Adam,  
1031 RMSprop, and evolutionary strategies [Soydaner, 2020] [Bingham et al., 2020], can  
1032 also be effective for training deep networks. The loss function serves as a crucial  
1033 component of the learning process in most optimization techniques, quantifying  
1034 how well the model's predictions match the actual targets. By providing a measure  
1035 of the model's performance, the loss function guides the chosen optimization  
1036 process, enabling the network to learn meaningful patterns from the training data.

### 1037 2.3.3 Challenges and Solutions

1038 Despite their power, deep learning systems face several significant challenges. One  
1039 of the most prominent issues is overfitting, where models may memorize training  
1040 data instead of learning generalizable patterns [Roelofs et al., 2019]. This challenge  
1041 is typically addressed through various regularization techniques such as dropout,  
1042 L1/L2 regularization, and early stopping [Roelofs et al., 2019]. Another critical  
1043 challenge is the substantial data requirements; deep learning models often need  
1044 massive amounts of training data to achieve good performance [Roelofs et al.,  
1045 2019], which can be a limiting factor in many applications. Additionally, the  
1046 complex, layered nature of deep learning models makes them difficult to interpret,  
1047 often referred to as "black boxes". This lack of transparency can be particularly  
1048 problematic in critical applications where understanding the decision-making  
1049 process is essential.

### 1050 2.3.4 Deep Learning Frameworks

1051 The development of powerful open-source frameworks has significantly  
1052 accelerated deep learning research and applications. TensorFlow, developed by

Google, provides a comprehensive ecosystem for building and deploying Machine Learning models [Abadi et al., 2016]. PyTorch, created by Facebook’s AI Research lab, offers dynamic computation graphs and has become particularly popular in research settings [Paszke et al., 2019]. Caffe, known for its speed and modularity, is widely used in computer vision applications [Jia et al., 2014]. These frameworks have democratized deep learning by providing efficient implementations of common operations, automatic differentiation for gradient computation, and GPU acceleration for faster training. They also offer pre-trained models and transfer learning capabilities, along with active communities for support and knowledge sharing. The combination of these frameworks with modern hardware has enabled researchers and practitioners to develop increasingly sophisticated models, pushing the boundaries of what’s possible in artificial intelligence. TensorFlow and PyTorch have emerged as the two most prominent frameworks in the deep learning landscape, as shown in Figure 2-6, which presents data from Google Trends over the last five years (as of April 2025).




Figure 2-6 Comparative Trends of the top two most popular Deep Learning Frameworks.

## 2.4 Datasets and Data Sources

Throughout the development of this work, multiple datasets were used for evaluation of ISS models. The common elements of all these datasets are that they

contain RGB images and are crowdsourced with multiple labelers, where not necessarily all labeler label all images.

As it has been mentioned in Chapter 1, the main goal of this work is mainly focused on crowdsourced histopathology images semantic segmentation, however, these datasets present the following challenges:

- Distribution of segmentation labels is not uniform across the image, since some tissues and structures are more common than others.
- Visualization of performance of the models in debug time (like per epoch analysis) is not simple for non experts in the subject, which makes it hard to evaluate whether the model is overfitting or not at a glance.

For these reasons, multiple datasets were created this work in the pursuit of an initial evaluation of performance of the models against more traditional and familiar images before the focus on histopathology images. Once a satisfactory performance was achieved, the focus was shifted to histopathology images and further tunings on the models were performed if needed.

In any case, both the emulated noisy annotations datasets and the histopathology datasets somehow contained ground truth aggregation, either from the original source (in the case of emulated noisy annotations), the expert annotation (if available) or from the aggregation of multiple labelers <sup>1</sup>.

### 2.4.1 Oxford-IIIT Pet Dataset

Using the techniques described later in Section 5.5.1, the Oxford-IIIT Pet Dataset [Parkhi et al., 2012] was used to create a dataset with emulated noisy annotations. This dataset is officially available in the ‘tensorflow\_datasets’ package, making it convenient to use in experiments. It is composed of 7349 images of pets, being 2371 of these images cats, and the rest 4978, dogs. Breed labels are included as well as segmentation masks for the classes: ‘background’, ‘border’, and ‘pet’. Figure 2-7 shows an example of the annotations in the original dataset.

---

<sup>1</sup>STAPLE in the case of histopathology datasets with no expert annotations available






Figure 2-7 Annotations in the Oxford-IIIT Pet data. From left to right: pet image, head bounding box, and trimap segmentation (blue: background region; red: ambiguous region; yellow: foreground region).

## 2.4.2 Real histopathology datasets

### Multi-Stain Breast Cancer Histological Dataset

The Multi-Stain Breast Cancer Histological Dataset [Weitz et al., 2023] represents one of the largest publicly available collections of whole slide images (WSIs) from surgical resection specimens of primary breast cancer patients. This dataset is particularly valuable for our work because it contains matched pairs of H&E and IHC-stained tissue sections from the same tumor, with a total of 4,212 WSIs from 1,153 patients. The IHC stains include important biomarkers such as ER (Estrogen Receptor), PGR (Progesterone Receptor), HER2 (Human Epidermal Growth Factor Receptor 2), and KI67 (Ki67 Proliferation Marker), which are routinely used in breast cancer diagnosis and treatment planning (more on staining techniques in Section 2.2.3).

The dataset's relevance to our work stems from several key aspects. The matched H&E and IHC stains allows for studying the consistency of segmentation across different staining modalities, which is crucial for understanding how different visualization methods affect annotation quality. With 1,153 patients, the dataset provides a robust foundation for training and evaluating segmentation models in a

1115 real-world clinical setting. The inclusion of routine diagnostic cases makes the  
 1116 dataset representative of actual clinical practice, where variations in staining  
 1117 quality and tissue preparation are common. Furthermore, the multiple biomarker  
 1118 stains (ER, PGR, HER2, KI67) enable to study how different tissue characteristics  
 1119 affect segmentation performance and annotator agreement.

1120 This dataset serves as an ideal testbed for our crowdsourced segmentation  
 1121 approach. It allows us to evaluate how different staining modalities affect  
 1122 annotator performance and agreement, while also providing insights into the  
 1123 relationship between tissue characteristics and segmentation difficulty. The  
 1124 dataset's comprehensive nature enables validation of our models' performance  
 1125 across different biomarker expressions and assessment of the generalizability of  
 1126 our approach to real-world clinical data.




Figure 2-8 Different staining techniques obtained from multi-stain breast cancer dataset [Weitz et al., 2023]. (a) shows H&E, (b) ER, (c) HER2, (d) Ki67 and (e) PGR. (f) shows an example of a WSI that was excluded since it contains multiple tissue sections.

## 1127 Structured Crowdsourcing Dataset for Histology Images

1128 The dataset presented by [Amgad et al., 2019] is particularly relevant to our work  
 1129 as it represents one of the first systematic studies of crowdsourced annotations in  
 1130 histopathology. The authors recruited 25 participants with varying levels of  
 1131 expertise (from senior pathologists to medical students) to delineate tissue regions  
 1132 in 151 breast cancer slides using the Digital Slide Archive platform, resulting in  
 1133 over 20,000 annotated tissue regions. The annotation efforts are focused on

1134 Triple-Negative Breast Cancer (TNBC), an aggressive genomic subtype that  
1135 comprises 15% of breast cancer cases [Plasilova et al., 2016].

1136 Key aspects of this dataset make it valuable for our work. The systematic  
1137 evaluation of inter-participant discordance revealed varying levels of agreement  
1138 across different tissue classes, with low discordance for tumor and stroma, and  
1139 higher discordance for more subjectively defined or rare tissue classes. The  
1140 inclusion of feedback from senior participants helped in curating high-quality  
1141 annotations, demonstrating that fully convolutional networks trained on these  
1142 crowdsourced annotations can achieve high accuracy (mean AUC=0.945). The  
1143 dataset also provides evidence that the scale of annotation data significantly  
1144 improves image classification accuracy.

1145 This dataset is particularly valuable for our work because it provides crucial  
1146 insights into how annotator expertise affects segmentation quality. It  
1147 demonstrates the feasibility of using crowdsourced annotations for training  
1148 accurate segmentation models, showing that even with varying levels of expertise,  
1149 aggregated annotations can produce reliable ground truth. The dataset includes a  
1150 systematic analysis of inter-annotator agreement, which is crucial for  
1151 understanding the challenges in crowdsourced histopathology segmentation and  
1152 informing the development of more robust segmentation approaches.




Figure 2-9 Screenshot of the DSA and HistomicsTK web interface while creating the crowdsourced annotations for the dataset presented by [Amgad et al., 2019].

1153

1154

1155

1156

CHAPTER

**THREE**

1157

## CHAINED GAUSSIAN PROCESSES

1158

### 3.1 Background and Related Methods

1159

1160

1161

1162

1163

In recent years, several approaches have been proposed to handle supervised learning problems in the context of multiple annotators [Yan et al., 2014] [Zhang, 2022]. This section reviews key methods that form the foundation for understanding chained Gaussian processes and their application to multiple annotator scenarios in regression and classification tasks.

1164

1165

#### 3.1.1 Kernel Alignment-Based Annotator Relevance Analysis

1166

1167

1168

The Kernel Alignment-Based Annotator Relevance Analysis (KAAR) method, introduced in [Gil-Gonzalez et al., 2018], addresses the challenge of estimating annotator expertise in scenarios where the ground truth is unavailable. The key

1169 innovation of **KAAR** lies in its use of **Centered Kernel Alignment (CKA)** to measure  
 1170 the similarity between input features and annotator labels.

1171 Given a dataset with input features  $\mathbf{X}$  and labels from multiple annotators  $\mathbf{Y}$ , **KAAR**  
 1172 computes a kernel matrix  $\mathbf{K}_\nu$  as a convex combination of  $R$  basis kernels:

$$\mathbf{K}_\nu = \sum_{r=1}^R \nu_r \mathbf{K}_r, \quad (3-1)$$

1173 where  $\mathbf{K}_r$  is computed using a kernel function over the input features, and  $\nu_r$   
 1174 represents the relevance weight for the  $r$ -th annotator. The weights are optimized  
 1175 by maximizing the CKA between  $\mathbf{K}_\nu$  and a target kernel  $\mathbf{F}$  computed from the  
 1176 input features:

$$\rho(\mathbf{K}_\nu, \mathbf{F}) = \frac{\langle \bar{\mathbf{K}}_\nu, \bar{\mathbf{F}} \rangle_{\text{F}}}{\|\bar{\mathbf{K}}_\nu\|_{\text{F}} \|\bar{\mathbf{F}}\|_{\text{F}}}, \quad (3-2)$$

1177 A key advantage of **KAAR** is its ability to capture dependencies between annotators  
 1178 through the kernel framework, without requiring explicit parametric modeling of  
 1179 these relationships.

### 1180 **3.1.2 Localized Kernel Alignment-Based Annotator** 1181 **Relevance Analysis**

1182 In [Gil-Gonzalez et al., 2021a] the authors extend **KAAR** by introducing locality in  
 1183 the kernel alignment framework. Rather than assuming constant annotator  
 1184 performance across the input space, **Localized Kernel Alignment-Based Annotator**  
 1185 **Relevance Analysis (LKAAR)** models the annotator relevance as a function of the  
 1186 input features. The kernel combination in **LKAAR** takes the form:

$$\mathbf{K}_q = \sum_{r=1}^R \mathbf{Q}_r \mathbf{K}_r \mathbf{Q}_r, \quad (3-3)$$

1187 where  $\mathbf{Q}_r$  is a diagonal matrix whose elements represent the local relevance of  
 1188 annotator  $r$ . These elements are computed as:

$$q_r(\mathbf{x}_n) = \begin{cases} \beta_0^{(r)} + \sum_{n'=1}^N \beta_{n'}^{(r)} \kappa_\beta(\mathbf{x}_n, \mathbf{x}_{n'}), & \text{if } n \in N_r \\ 0, & \text{Otherwise} \end{cases} \quad (3-4)$$

1189 This localized approach allows **LKAAR** to better model inconsistent annotators  
 1190 whose performance varies across different regions of the input space [**Gil-Gonzalez**  
 1191 **et al., 2021a**].

### 1192 3.1.3 Chained Gaussian Processes Model

1193 The **CGP** model provides a probabilistic framework for modeling multiple outputs  
 1194 through a chain of Gaussian processes [**Saul et al., 2016**]. In the context of multiple  
 1195 annotators, **CGP** can be used to model both the ground truth and annotator  
 1196 behavior. The key idea is to model a likelihood function with  $J$  parameters as in  
 1197 [**Giraldo et al., 2022**]:

$$p(\mathbf{y}|\mathbf{X}, \boldsymbol{\theta}) = \prod_{n=1}^N p(y_n | \theta_1(\mathbf{x}_n), \dots, \theta_J(\mathbf{x}_n)), \quad (3-5)$$

1198 where each parameter  $\theta_j(\mathbf{x})$  is modeled using a separate Gaussian process. This  
 1199 framework provides remarkable flexibility in modeling complex relationships  
 1200 within the data [**Giraldo et al., 2022**]. The use of Gaussian processes enables  
 1201 natural uncertainty quantification through posterior distributions, allowing the  
 1202 model to express varying degrees of confidence in its predictions. The framework

1203 inherently handles missing data through its probabilistic nature, and the automatic  
 1204 relevance determination property of GPs helps identify the most important  
 1205 features for prediction [Giraldo et al., 2022].

1206 The CGP framework becomes particularly powerful when integrated with concepts  
 1207 from KAAR and LKAAR. By incorporating kernel alignment principles, the model  
 1208 can make more informed choices about kernel functions, better capturing the  
 1209 underlying structure of the data [Giraldo et al., 2022]. The framework can be  
 1210 extended to include input-dependent lengthscales, allowing for varying degrees of  
 1211 smoothness across the input space [Giraldo et al., 2022]. Furthermore, through  
 1212 multi-output Gaussian Process (GP) formulations, the model can capture complex  
 1213 dependencies capture complex dependencies between different annotators,  
 1214 providing a rich representation of how different experts' opinions relate to each  
 1215 other.

1216 However, implementing CGPs in practice presents several significant challenges.  
 1217 The computational complexity of the model scales cubically with the dataset size,  
 1218 making it potentially prohibitive for large-scale applications [Liu et al., 2020]. The  
 1219 need for approximate inference techniques in non-Gaussian likelihood scenarios  
 1220 can introduce additional complexity and potential approximation errors  
 1221 [Rasmussen and Williams, 2006a]. When dealing with high-dimensional inputs,  
 1222 the model may struggle with the curse of dimensionality, requiring careful feature  
 1223 selection or dimensionality reduction preprocessing [Rasmussen and Williams,  
 1224 2006a]. Additionally, the selection of appropriate kernel functions remains a  
 1225 critical challenge, as it significantly impacts the model's performance and its ability  
 1226 to capture relevant patterns in the data [Rasmussen and Williams, 2006b].

### 1227 3.1.4 Regularized Chained Deep Neural Network

1228 In [Gil-Gonzalez et al., 2021b], the authors propose a deep learning approach to  
 1229 multiple annotator learning, inspired by the chained Gaussian processes model. The



method models both the ground truth estimation and annotator reliability through a deep neural network architecture. The likelihood function in **Regularized Chained Deep Neural Network (RCDNN)** takes the form:

$$p(\mathbf{Y}|\boldsymbol{\theta}) = \prod_{n=1}^N \prod_{r \in R_n} \left( \prod_{k=1}^K \zeta_{n,k}^{\delta(y_n^{(r)} - k)} \right)^{\lambda_n^{(r)}} \left( \frac{1}{K} \right)^{1 - \lambda_n^{(r)}}, \quad (3-6)$$

where  $\lambda_n^{(r)}$  represents the reliability of annotator  $r$  for instance  $n$ , and  $\zeta_{n,k}$  is the estimated probability of class  $k$  for instance  $n$ .

The **RCDNN** architecture is built upon a sophisticated multi-layer structure that combines several key innovations in deep learning [Gil-Gonzalez et al., 2021b]. At its core, it employs multiple dense layers with carefully chosen activation functions, each designed to capture different aspects of the annotator-data relationship. To prevent overfitting, a comprehensive regularization strategy is implemented, combining both L1 and L2 norms with dropout mechanisms [Gil-Gonzalez et al., 2021b]. The network architecture splits into specialized branches: one dedicated to estimating the ground truth labels, and another focused on modeling annotator reliability. This dual-branch approach allows the network to simultaneously learn both the true underlying patterns in the data and the individual characteristics of each annotator. Furthermore, the implementation of Monte Carlo dropout during prediction provides a mechanism for uncertainty estimation, offering not just predictions but also confidence measures for those predictions.

## 3.2 Tooling and Implementation

The implementation of Chained Gaussian Processes for multiple annotator learning requires robust and flexible software tools that can handle both the probabilistic

1252 nature of Gaussian processes and the complexity of deep architectures [Saul et al.,  
1253 2016]. In this work, two key frameworks are leveraged: GPflow and GPflux. These  
1254 tools provide the foundation for building scalable and efficient implementations of  
1255 our models while maintaining the theoretical rigor necessary for proper uncertainty  
1256 quantification.

1257 The choice of these frameworks is motivated by several factors. First, both are  
1258 built on top of TensorFlow, providing access to automatic differentiation and GPU  
1259 acceleration capabilities. Second, they offer a modular design that allows for easy  
1260 extension and modification, which is crucial when implementing novel model  
1261 architectures. Third, they provide robust implementations of various Gaussian  
1262 process models and inference methods, allowing us to focus on the specific  
1263 challenges of multiple annotator learning rather than reimplementing basic GP  
1264 functionality.

### 1265 3.2.1 GPflow Overview

1266 GPflow is a Python package for building Gaussian process models in TensorFlow.  
1267 Developed as a modern and scalable alternative to GPy, it leverages TensorFlow’s  
1268 computational capabilities to provide efficient implementation of Gaussian process  
1269 models. The framework is designed with four key principles in mind: ease of use,  
1270 computational efficiency, extensibility, and automatic differentiation. This work  
1271 was originally published in [de G. Matthews et al., 2016].

1272 At its core, GPflow provides implementations of various Gaussian process models  
1273 including Gaussian Process Regression (GPR), Gaussian Process Monte Carlo  
1274 (GPMC), Sparse Gaussian Process Regression (SGPR), and Sparse Gaussian Process  
1275 Monte Carlo (SGPMC) [de G. Matthews et al., 2016]. These implementations are  
1276 built on top of TensorFlow’s computational graph framework, allowing for  
1277 automatic differentiation and GPU acceleration where available.

1278 The framework’s architecture is modular, with clear separation between model  
1279 components such as kernels, likelihoods, and optimization routines. This  
1280 modularity makes it straightforward to extend existing models or create new ones.  
1281 GPflow includes a comprehensive suite of kernel functions, ranging from the basic  
1282 RBF (Radial Basis Function) kernel to more sophisticated ones like the Spectral  
1283 Mixture kernel.

1284 One of GPflow’s distinguishing features is its robust handling of model parameters.  
1285 It provides facilities for parameter transformation (ensuring parameters remain in  
1286 valid ranges during optimization), prior specification, and flexible optimization  
1287 strategies. The framework also includes utilities for model evaluation and  
1288 prediction, with built-in support for uncertainty estimation.

## 1289 3.2.2 GPflux Framework

1290 GPflux extends GPflow to enable the construction and training of Deep Gaussian  
1291 Process (DGP) models using modern deep learning practices. Built on top of both  
1292 GPflow and TensorFlow, GPflux provides a flexible framework for implementing  
1293 deep probabilistic models that combine the expressiveness of deep neural  
1294 networks with the uncertainty quantification capabilities of Gaussian processes.  
1295 This work was originally published in [Dutordoir et al., 2021].

1296 The framework implements Deep Gaussian Processes as a composition of Gaussian  
1297 process layers, where each layer transforms its inputs through a Gaussian process  
1298 mapping. This hierarchical structure allows the model to capture complex,  
1299 non-stationary patterns in data while maintaining uncertainty estimates  
1300 throughout the entire processing pipeline. GPflux handles this complexity through  
1301 a carefully designed architecture that separates the concerns of model  
1302 specification, inference, and prediction.

1303 A key innovation in GPflux is its integration with Keras-style layers, allowing  
1304 practitioners to combine deterministic neural network layers with Gaussian

process layers in a single model. This hybrid approach enables the construction of models that can leverage both the structural assumptions encoded in neural network architectures and the flexibility of Gaussian processes.

For inference, GPflux implements various approximate methods, including stochastic variational inference with inducing points, which makes it possible to train deep Gaussian process models on large datasets. The framework provides built-in support for minibatch training, making it practical to work with datasets that don't fit in memory.

### 3.3 Experimental Setup

Synthetic datasets were used to evaluate the proposed methods in classification and regression tasks, following very similar procedures to the ones used in [Gil-Gonzalez et al., 2023], which are as follows:

#### 3.3.1 Classification

A fully synthetic data for a 1-D ( $P = 1$ ) multiclass classification problem with  $K = 3$  classes is generated. The input feature matrix  $\mathbf{X}$  is constructed by randomly sampling  $N = 100$  points from a uniform distribution over the interval  $[0, 1]$ . For each sample  $n$ , the true label is determined by computing three target values:

$$\begin{aligned} t_{n,1} &= \sin(2\pi x_n), \\ t_{n,2} &= -\sin(2\pi x_n), \\ t_{n,3} &= -\sin(2\pi(x_n + 0.25)) + 0.5, \end{aligned} \tag{3-7}$$

and assigning the class label as the index  $i \in \{1, 2, 3\}$  that maximizes  $t_{n,i}$ . This construction ensures a nontrivial, nonlinear decision boundary between the classes [Gil-González and Álvarez Meza, 2023].

For evaluation, the test set is generated by extracting 200 equally spaced samples from the interval  $[0, 1]$ , providing a dense coverage of the input space to assess generalization performance.

Note that at this point, the fully synthetic dataset does not contain real annotator labels. Therefore, it is necessary to simulate annotator labels as corrupted versions of the hidden ground truth. In these experiments, the simulation is designed to account for: 1) dependencies among annotators, and 2) labeler performance that varies as a function of the input features.

To achieve this, an **SLFM**-based approach (termed SLFM-C) is employed, which generates annotator labels as follows:

1. **Definition of deterministic functions:** Define  $Q$  deterministic functions  $\hat{\mu}_q : \mathcal{X} \rightarrow \mathbb{R}$  and their combination parameters  $\hat{w}_{l,r,q} \in \mathbb{R}$ , for all  $r \in R, n \in N$ .
2. **Annotator performance computation:** Compute  $\hat{f}_{l,r,n} = \sum_{q=1}^Q \hat{w}_{l,r,q} \hat{\mu}_q(\hat{x}_n)$ , where  $\hat{x}_n \in \mathbb{R}$  is the  $n$ th component of  $\hat{\mathbf{x}} \in \mathbb{R}^N$ . Here,  $\hat{\mathbf{x}}$  is the 1-D representation of the input features in  $\mathbf{X}$ , obtained using the t-distributed Stochastic Neighbor Embedding (t-SNE) approach.
3. **Reliability calculation:** Calculate  $\hat{\lambda}_n^r = \varsigma(\hat{f}_{l,r,n})$ , where  $\varsigma(\cdot) \in [0, 1]$  is the inverse exponential sigmoid function.

4. **Label assignment:** The  $r$ th annotator's label for the  $n$ th sample is given by

$$y_n^r = \begin{cases} y_n, & \text{if } \hat{\lambda}_n^r \geq 0.5 \\ \tilde{y}_n, & \text{if } \hat{\lambda}_n^r < 0.5 \end{cases}$$

where  $y_n$  is the true label and  $\tilde{y}_n$  is a flipped version of the actual label  $y_n$ .

This procedure ensures that the simulated annotator labels exhibit both inter-annotator dependencies and input-dependent reliability.

### 1347 3.3.2 Regression

1348 A fully synthetic dataset is also generated for a 1-D regression problem, where the  
1349 ground truth for the  $n$ th sample is defined as:

$$y_n = \sin(2\pi x_n) \sin(6\pi x_n) \quad (3-8)$$

1350 Here, the input matrix  $\mathbf{X}$  is constructed by randomly sampling 100 points uniformly  
1351 from the interval  $[0, 1]$ . The test set is created by extracting equally spaced samples  
1352 from the same interval, ensuring a clear separation between training and testing  
1353 data.

1354 To simulate the presence of multiple annotators, the ground truth is corrupted with  
1355 annotator-specific noise. For each annotator  $r$ , the observed label for the  $n$ th sample  
1356 is given by:

$$y_n^r = y_n + \epsilon_n^r$$

1357 where  $\epsilon_n^r \sim \mathcal{N}(0, v_n^r)$  is Gaussian noise, and  $v_n^r$  represents the error variance for  
1358 the  $r$ th annotator, which may depend on the input features. This setup allows us  
1359 to model both the bias and the reliability of each annotator, as well as potential  
1360 dependencies between annotators.

1361 To further increase realism, the error variance  $v_n^r$  can be modeled as a function of  
1362 the input features, capturing scenarios where annotator reliability varies across the  
1363 input space, thus:

$$\epsilon_n^r \sim \mathcal{N}(0, v_n^r)$$

1364 being  $v_n^r$ , the  $r$ th annotator's error variance for the  $n$ th sample.

1365 Again, for modeling the error as a function of the input features, an SLFM-based  
1366 approach (termed SLFM-R) is employed to build the noisy labels, as follows:

- 1367 **1. Definition of deterministic functions:** Define  $Q$  deterministic functions  $\hat{\mu}_q :$   
 1368  $\mathcal{X} \rightarrow \mathbb{R}$  and their combination parameters  $\hat{w}_{l,r,q} \in \mathbb{R}$ , for all  $r \in R, n \in N$ .
- 1369 **2. Annotator performance computation:** Compute  $\hat{f}_{l,r,n} = \sum_{q=1}^Q \hat{w}_{l,r,q} \hat{\mu}_q(\hat{x}_n)$ ,  
 1370 where  $\hat{x}_n \in \mathbb{R}$  is the  $n$ th component of  $\hat{\mathbf{x}} \in \mathbb{R}^N$ . Here,  $\hat{\mathbf{x}}$  is the 1-D  
 1371 representation of the input features in  $\mathbf{X}$ , obtained using the t-distributed  
 1372 Stochastic Neighbor Embedding (t-SNE) approach.
- 1373 **3. Compute variance :** Finally, determine  $v_n^r = \exp(\hat{f}_{l,r,n})$ .

## 1374 3.4 Results


1375 The built synthetic datasets are used to train the **CCGPMA** model. Emulated  
 1376 classification labels following the true label determining function from 3-7 can be  
 1377 seen in Figure 3-1. On the other hand, the ground truth and noisy labels emulation  
 1378 for annotators across the feature space for regression task can be seen in Figure  
 1379 3-3.

1380 Using both GPflux and GPflow for the regression and classifications tasks, leads to  
 1381 the following combinations experiments: GPflux-CCGPMA Regression,  
 1382 GPflow-CCGPMA Regression, GPflux-CCGPMA Binary Classification,  
 1383 GPflow-CCGPMA Binary Classification, GPflow-CCGPMA Multiclass Classification,  
 1384 GPflux-CCGPMA Multiclass Classification. All of the experiments converged  
 1385 successfully with almost identical results and close to no difference in  
 1386 performance.


## 1387 3.5 Conclusion

### 1388 3.5.1 Comparative Analysis

1389 The evolution of multiple annotator learning methods represents a fascinating  
 1390 progression in how we approach the challenge of learning from diverse expert



(a) Ground truth labels across the feature space



(b) Noisy emulations for annotators across the feature space for classification task

Figure 3-1 Synthetic dataset overview for classification task






Figure 3-2 Classification results for the CCGPMA model using GPflux.





Figure 3-3 Ground truth and noisy labels emulation for annotators across the feature space for regression task.



(a) Trained CCGPMA predictions on the test set, for inference of both the hidden ground truth and the annotators' reliabilities across the feature space.

Figure 3-4 Regression results for the CCGPMA model using GPflux.

opinions. **KAAR** established a foundation with its non-parametric approach through kernel alignment, though it operates under the assumption of constant annotator performance across the input space. **LKAAR** advanced this concept by introducing input-dependent annotator behavior, recognizing that expert performance often varies across different types of inputs. **RCDNN** brought the power and flexibility of deep learning to the problem, offering excellent scalability but requiring careful attention to regularization and architecture design. **CCGPMA** represents perhaps the most theoretically elegant approach, providing a principled probabilistic framework, though it faces significant computational challenges in practical applications.

The selection of an appropriate method for a given application requires careful consideration of several factors. The size and dimensionality of the dataset play a crucial role, as some methods scale better than others with increasing data volume. The need for uncertainty quantification may favor probabilistic approaches like CGP, while computational constraints might push towards more efficient methods like KAAR. The presence of input-dependent annotator behavior could make **LKAAR** or **RCDNN** more appropriate choices. Additionally, when interpretability is a key requirement, as in many medical applications, methods like **KAAR** and **LKAAR** might be preferred over the more complex Deep Learning approaches. Table 3-1 provides a comprehensive comparison of the key characteristics and capabilities of each method discussed.

Table 3-1 Comparison of Multiple Annotator Learning Methods

| Characteristic                     | <b>KAAR</b><br>[Gil-Gonzalez et al., 2018] | <b>LKAAR</b><br>[Gil-Gonzalez et al., 2021a] | <b>RCDNN</b><br>[Gil-Gonzalez et al., 2021b] | <b>CGP</b><br>[Gil-González and Álvarez Meza, 2023] |
|------------------------------------|--------------------------------------------|----------------------------------------------|----------------------------------------------|-----------------------------------------------------|
| Input-dependent annotator behavior | No                                         | Yes                                          | Yes                                          | Yes                                                 |
| Annotator dependencies modeling    | Yes                                        | Yes                                          | Yes                                          | Yes                                                 |
| Uncertainty quantification         | No                                         | No                                           | Partial                                      | Yes                                                 |
| Computational scalability          | High                                       | Medium                                       | High                                         | Low                                                 |
| Missing data handling              | Yes                                        | Yes                                          | Yes                                          | Yes                                                 |
| Non-linear relationships           | Yes                                        | Yes                                          | Yes                                          | Yes                                                 |
| Training complexity                | Low                                        | Medium                                       | High                                         | High                                                |
| Hyperparameter sensitivity         | Low                                        | Medium                                       | High                                         | Medium                                              |
| Interpretability                   | High                                       | Medium                                       | Low                                          | Medium                                              |
| Memory requirements                | Low                                        | Medium                                       | High                                         | High                                                |

**Note:** Partial uncertainty quantification in **RCDNN** refers to the use of Monte Carlo dropout for approximate uncertainty estimation.

### 1412 3.5.2 Implementation Considerations

1413 When implementing Chained Gaussian Processes for multiple annotator learning  
 1414 using these tools, several key considerations must be addressed. First, the model  
 1415 architecture must be carefully designed to capture both the individual annotator  
 1416 characteristics and their interdependencies [Gil-González and Álvarez Meza, 2023].  
 1417 This requires extending the basic GP layers provided by GPflux to handle multiple  
 1418 output streams - one for each annotator’s reliability model.

1419 Second, the inference procedure must be adapted to handle the unique challenges  
 1420 of multiple annotator data, such as missing labels and varying levels of annotator  
 1421 expertise across different regions of the input space. The variational inference  
 1422 capabilities of GPflux provide a solid foundation for this, but careful attention  
 1423 must be paid to the design of the variational approximation to ensure it captures  
 1424 the relevant uncertainty in both the ground truth and annotator reliability  
 1425 estimates.

1426 Third, kernel selection and design presents a critical challenge in implementing the  
 1427 CGP model. The choice of kernel functions significantly impacts the model’s ability  
 1428 to capture the underlying structure of the data and the relationships between  
 1429 annotators [Rasmussen and Williams, 2006b]. Several considerations must be  
 1430 taken into account:

- 1431 • **Kernel Complexity:** The model requires multiple kernel functions to capture  
 1432 different aspects of the data - one for the ground truth and additional kernels  
 1433 for each annotator’s reliability model. This multi-kernel setup increases the  
 1434 complexity of kernel selection and optimization.
- 1435 • **Input-Dependent Behavior:** Since annotator performance often varies across  
 1436 the input space, the kernel functions must be flexible enough to capture this  
 1437 non-stationary behavior. This may require using composite kernels or kernels  
 1438 with input-dependent lengthscales.

- 1439 • **Inter-Annotator Dependencies:** The kernel design must account for potential  
1440 correlations between different annotators' behaviors. This can be achieved  
1441 through careful design of the kernel structure in the SLFM framework, where  
1442 the combination of basis kernels must be chosen to effectively model these  
1443 dependencies.
- 1444 • **Computational Efficiency:** The choice of kernel functions also impacts  
1445 computational performance. While more expressive kernels may better  
1446 capture complex patterns, they often come with increased computational  
1447 cost. This trade-off must be carefully considered, especially when dealing  
1448 with large datasets.

1449 Finally, computational efficiency becomes a critical concern when scaling to large  
1450 datasets with multiple annotators. The minibatch training capabilities of both  
1451 frameworks help address this, but additional considerations such as appropriate  
1452 inducing point selection and batch size tuning become important for achieving  
1453 good performance in practice.

### 1454 3.5.3 Final Remarks

1455 In this chapter, we have presented a comprehensive overview of the Chained  
1456 Gaussian Processes model and its application to multiple annotator learning. We  
1457 have discussed the theoretical underpinnings of the model, its implementation  
1458 using GPflow and GPflux, and its evaluation on synthetic datasets. It was shown in  
1459 [Gil-Gonzalez et al., 2023] that CCGPMA outperforms other method in state of the  
1460 art methods in terms of accuracy and uncertainty quantification (see 1.3.1).

1461 While Gaussian Processes have demonstrated remarkable success in regression  
1462 and classification tasks, their application to image segmentation presents  
1463 significant computational challenges [Rasmussen and Williams, 2006b]. The cubic  
1464 computational complexity of GPs ( $\mathcal{O}(N^3)$ ), which scales with the number of data

1465 points, becomes particularly problematic when dealing with high-resolution  
1466 images where each pixel can be considered a data point [Liu et al., 2020]. This  
1467 limitation makes traditional GP approaches impractical for large-scale image  
1468 segmentation tasks, where the number of pixels can easily reach into the millions.  
1469 Alternative approaches, such as deep learning-based segmentation methods, often  
1470 provide more computationally efficient solutions for these high-dimensional  
1471 problems, though at the cost of uncertainty quantification capabilities that GPs  
1472 naturally provide [Kendall and Gal, 2017] [Wilson et al., 2016].

1473

1474

1475

1476

---

CHAPTER

**FOUR**

---

1477

1478

## TRUNCATED GENERALIZED CROSS ENTROPY FOR SEGMENTATION

1479

### **4.1 Loss Functions for Multiple Annotators**

1480

1481

1482

1483

1484

1485

1486

As mentioned in Section 2.3.2, a loss function is a key element for defining the objective function of a deep learning model. The categorical cross-entropy loss is a common loss function for classification tasks. However, in the case of multiple annotators, the categorical cross-entropy loss is not able to handle the varying reliability of the annotators. In this section, we will propose a loss function that is able to handle multiple annotators' segmentation masks while accounting for their varying reliability across different regions of the image.

### 1487 4.1.1 Generalized Cross Entropy

1488 The Generalized Cross Entropy (GCE) loss function was first introduced by [Zhang  
1489 and Sabuncu, 2018] as a robust alternative to the standard cross-entropy loss,  
1490 particularly effective in handling noisy labels. Let us first consider the Cross  
1491 Entropy (CE) and Mean Absolute Error (MAE) loss functions:

$$MAE(\mathbf{y}, f(\mathbf{x})) = \|\mathbf{y} - f(\mathbf{x})\|_1 \quad (4-1)$$

$$CE(\mathbf{y}, f(\mathbf{x})) = \sum_{k=1}^K y_k \log(f_k(\mathbf{x})) \quad (4-2)$$

1492 where  $y_k \in \mathbf{y}$ ,  $f_k(\mathbf{x}) \in f(\mathbf{x})$ , and  $\|\cdot\|_1$  stands for the  $l_1$ -norm. Of note,  $\mathbf{1}^\top \mathbf{y} =$   
1493  $\mathbf{1}^\top f(\mathbf{x}) = 1$ ,  $\mathbf{1} \in \{1\}^K$  being an all-ones vector. In addition, the MAE loss can be  
1494 rewritten for softmax outputs, yielding:

$$MAE(\mathbf{y}, f(\mathbf{x})) = 2(1 - \mathbf{y}^\top \odot f(\mathbf{x})) \quad (4-3)$$

1495 where  $\odot$  stands for the element-wise product.

1496 The CE loss function exhibits several distinct characteristics that make it  
1497 particularly sensitive to noisy labels. It is unbounded from above and heavily  
1498 penalizes confident but wrong predictions, which can lead to instability in the  
1499 presence of label noise. In contrast, the MAE loss offers a more robust alternative  
1500 with its bounded nature and symmetric properties in softmax-based  
1501 representations. While the MAE loss is more resilient to noisy labels, it tends to  
1502 train more slowly due to its equal weighting of all mistakes regardless of  
1503 confidence.

1504 The GCE loss function, as defined by the authors in [Zhang and Sabuncu, 2018],  
1505 provides a flexible framework that bridges these two approaches:



$$GCE(\mathbf{y}, f(\mathbf{x})) = 2 \frac{1 - (\mathbf{1}^\top (\mathbf{y} \odot f(\mathbf{x})))^q}{q}, \quad (4-4)$$

1506 with  $q \in (0, 1]$ . Remarkably, the limiting case for  $q \rightarrow 0$  in GCE is equivalent to the  
 1507 CE expression, and when  $q = 1$ , it equals the MAE loss. In addition, the GCE holds  
 1508 the following gradient with regard to  $\theta$ :

$$\frac{\partial GCE(\mathbf{y}, f(\mathbf{x}; \theta)|k)}{\partial \theta} = -f_k(\mathbf{x}; \theta)^{q-1} \nabla_\theta f_k(\mathbf{x}; \theta). \quad (4-5)$$

1509 The GCE loss combines the best of both worlds, offering robustness to label noise  
 1510 while maintaining the convexity property necessary for optimization. The  
 1511 truncation parameter  $q$  provides a mechanism to control the sensitivity to outliers,  
 1512 allowing for fine-tuning of the loss function's behavior based on the specific  
 1513 characteristics of the dataset.

## 1514 4.1.2 Extension to Multiple Annotators

1515 In the context of multiple annotators, we need to consider the varying reliability  
 1516 of each annotator across different regions of the image. Let's consider a  $k$ -class  
 1517 multiple annotators segmentation problem with the following data representation:

$$\mathbf{X} \in \mathbb{R}^{W \times H}, \{\mathbf{Y}_r \in \{0, 1\}^{W \times H \times K}\}_{r=1}^R; \hat{\mathbf{Y}} \in [0, 1]^{W \times H \times K} = \mathbf{f}(\mathbf{X}) \quad (4-6)$$

1518 where the segmentation mask function maps the input to output as:

$$\mathbf{f}: \mathbb{R}^{W \times H} \rightarrow [0, 1]^{W \times H \times K} \quad (4-7)$$

1519 The segmentation masks  $\mathbf{Y}_r$  satisfy the following condition for being a softmax-like  
 1520 representation:

$$\mathbf{Y}_r[w, h, :] \mathbf{1}_k^\top = 1; \quad w \in W, h \in H \quad (4-8)$$

### 1521 4.1.3 Reliability Maps and Truncated GCE

1522 The key innovation in our approach is the introduction of reliability maps  $\Lambda_r$  for  
1523 each annotator:

$$\left\{ \Lambda_r(\mathbf{X}; \theta) \in [0, 1]^{W \times H} \right\}_{r=1}^R \quad (4-9)$$

1524 These reliability maps serve as a mechanism to estimate the confidence of each  
1525 annotator at every spatial location  $(w, h)$  in the image. By learning these maps  
1526 jointly with the segmentation model, the network gains the ability to adapt to the  
1527 varying levels of expertise across different regions of the image. This approach  
1528 allows for dynamic weighting of each annotator's contribution based on their  
1529 estimated reliability in specific areas, effectively handling cases where annotators  
1530 might demonstrate different levels of expertise in different parts of the image.

1531 The proposed Truncated Generalized Cross Entropy for Semantic Segmentation  
1532 (TGCE<sub>SS</sub>) combines the robustness of GCE with the flexibility of reliability maps.  
1533 Let us first define the key components:

$$\mathcal{L}_{GCE}(\mathbf{Y}_r, f(\mathbf{X}; \theta)) = \mathbb{E}_k \left\{ \mathbf{Y}_r \odot \left( \frac{\mathbf{1}_{W \times H \times K} - f(\mathbf{X}; \theta)^{\odot q}}{q} \right) \right\} \quad (4-10)$$

$$\mathcal{L}_{prior} = \frac{\mathbf{1}_{W \times H} - \left( \frac{1}{K} \mathbf{1}_{W \times H} \right)^{\odot q}}{q} \quad (4-11)$$

1534 where  $\mathcal{L}_{GCE}$  represents the GCE loss component and  $\mathcal{L}_{prior}$  represents the uniform  
1535 prior component. The complete TGCE<sub>SS</sub> loss function is then defined as:

$$\begin{aligned} TGCE_{SS}(\mathbf{Y}_r, f(\mathbf{X}; \theta) | \mathbf{r}(\mathbf{X}; \theta)) = \mathbb{E}_{r,w,h} \left\{ \right. \\ \Lambda_r(\mathbf{X}; \theta) \odot \mathcal{L}_{GCE} + \tilde{C}(\mathbf{1}_{W \times H} - \Lambda_r(\mathbf{X}; \theta)) \odot \mathcal{L}_{prior}; \\ \left. w \in W, h \in H; r \in R \right\} \end{aligned} \quad (4-12)$$

where  $q \in (0, 1)$  controls the MAE or CE level in the same way as in the GCE loss function.  $\tilde{C}$  controls is an additional hyperparameter which controls the tolerance to unreliable annotations. The loss function consists of two main components:

- The first term weighted by  $\Lambda_r$  represents the GCE loss for regions where the annotator is considered reliable
- The second term weighted by  $(1 - \Lambda_r)$  and  $\tilde{C}$  provides a uniform prior for regions where the annotator is considered unreliable

A summary of the loss function mechanism is shown in Figure 4-1.




Figure 4-1 Working mechanism of the proposed Truncated Generalized Cross Entropy for Semantic Segmentation (TGCE<sub>SS</sub>) loss function. The loss combines reliability maps  $\Lambda_r$  with the model's predictions  $f(\mathbf{X}; \theta)$  to compute a weighted loss that accounts for annotator reliability across different image regions.

## 4.2 Chapter overview

This chapter presents a novel approach to handle multiple annotators in semantic segmentation tasks through the introduction of the Truncated Generalized Cross Entropy for Semantic Segmentation (TGCE<sub>SS</sub>) loss function. The key contributions and novelties of this work can be summarized as follows:

- **Reliability-aware Learning:** Unlike traditional approaches that treat all annotators equally, the proposed method introduces reliability maps ( $\Lambda_r$ ) that dynamically estimate the confidence of each annotator at every spatial location in the image. This allows the model to adapt to varying levels of expertise across different regions.
- **Robust Loss Function:** The proposed  $TGCE_{SS}$  combines the benefits of both Cross Entropy and Mean Absolute Error losses through the Generalized Cross Entropy framework. The truncation parameter  $q$  provides a flexible mechanism to control the sensitivity to noisy labels while maintaining the convexity property necessary for optimization.
- **Unified Framework:** The approach provides a unified framework that handles both reliable and unreliable regions through a weighted combination of GCE loss and uniform prior. This is particularly important in medical imaging where annotator expertise can vary significantly across different anatomical structures.

The main novelty of this approach lies in its ability to simultaneously learn the segmentation task while estimating annotator reliability in a spatially-aware manner. This represents a significant advancement over existing methods that either:

- Treat all annotators equally without considering their varying expertise
- Use global reliability scores that don't account for spatial variations in annotator performance
- Rely on fixed, pre-defined reliability measures rather than learning them from data

The proposed  $TGCE_{SS}$  loss function represents a step forward in handling the complex nature of multiple annotator scenarios in semantic segmentation, particularly in medical imaging applications where expert knowledge can vary significantly across different regions of interest. This work provides a foundation for more sophisticated approaches to handling multiple annotators in deep learning-based segmentation tasks.

1579

1580

1581

1582

---

CHAPTER

**FIVE**

---

1583

1584

## CHAINED DEEP LEARNING AND $TGCE_{SS}$ LOSS FOR IMAGE SEGMENTATION

1585

1586

1587

1588

As mentioned in Chapter 1, U-shaped Convolutional Neural Networks (CNNs) have been proven a good solution for Semantic segmentation (SS) tasks in medical images, due to their ability to capture both global and local information with small datasets.

1589

1590

1591

1592

1593

The proposed model architecture combines the strengths of U-Net with a ResNet-34 backbone, specifically designed to work with the  $TGCE_{SS}$  loss function by providing a multi-output network which focuses on two tasks: **Multiclass Semantic Segmentation** and **Reliability Estimation**. The architecture is illustrated in Figure 5-2.




Figure 5-1 Original U-NET architecture from [Ronneberger et al., 2015].

## 5.1 Backbone Architecture

The model employs a pre-trained ResNet-34 as its encoder backbone, leveraging its deep residual learning framework for efficient feature extraction. The choice of ResNet-34 provides several key advantages: efficient feature extraction through residual connections, pre-trained weights that capture rich visual representations, and stable gradient flow during training [He et al., 2015]. The ResNet-34 backbone is modified to serve as the encoder in our UNET architecture by removing the final fully connected layer and utilizing the feature maps from different stages of the network for skip connections.

### 5.1.1 U-Net Architecture

The U-Net architecture follows a traditional encoder-decoder structure with skip connections, where the encoder path implements the ResNet-34 structure. The




Figure 5-2 Overview of the proposed model architecture.

decoder path employs transposed convolutions for upsampling, creating a symmetrical architecture that effectively captures both high-level and low-level features. The architecture incorporates four downsampling stages in the encoder, corresponding to the ResNet-34 blocks, and four upsampling stages in the decoder. These stages are connected through skip connections that bridge corresponding encoder and decoder stages, allowing the network to preserve fine-grained details. Each convolution operation is followed by batch normalization and ReLU activation to ensure stable training and effective feature learning.

This U-Net architecture is expected to be coupled with a reliability branch which leads to a multi-output network. Three different approaches are considered for this branch: **scalar**, **features**, and **pixel**.

## 5.2 Reliability Branch

### 5.2.1 Scalar Reliability Estimation

The scalar reliability estimation branch is the simplest one, in which the reliability branch is a parallel branch to the segmentation branch, which leads to a

1621 multi-output network. The final output consists of  $R$  reliability maps  $\Lambda_r$ , one for  
 1622 each annotator, with values constrained to the  $[0, 1]$  range through a sigmoid  
 1623 activation function. This design allows the network to learn and adapt to the  
 1624 varying reliability of different annotators across different regions of the image.

1625 The scalar reliability branch is implemented by taking the bottleneck features from  
 1626 the encoder and applying a global average pooling operation followed by a dense  
 1627 layer with sigmoid activation. This results in a single reliability value per annotator  
 1628 for each input image, making it computationally efficient but less spatially aware of  
 1629 reliability variations within the image.

## 1630 5.2.2 Feature-based Reliability Estimation

1631 The feature-based reliability estimation approach provides a more detailed spatial  
 1632 understanding of annotator reliability by operating on the bottleneck features of  
 1633 the encoder. This approach maintains the spatial information from the encoder's  
 1634 bottleneck while being more computationally efficient than pixel-wise reliability  
 1635 estimation.

1636 The feature-based reliability branch is implemented by applying two convolutional  
 1637 layers with batch normalization to the bottleneck features:

- 1638 • A first convolutional layer with 32 filters to process the bottleneck features
- 1639 • A second convolutional layer that produces  $R$  feature maps, one for each  
 1640 annotator

1641 This approach provides a coarse-grained spatial reliability map that captures  
 1642 regional variations in annotator reliability while maintaining computational  
 1643 efficiency. The resulting reliability maps have the same spatial dimensions as the  
 1644 bottleneck features, providing a middle ground between scalar and pixel-wise  
 1645 reliability estimation.



### 5.2.3 Pixel-wise Reliability Estimation

The pixel-wise reliability estimation approach provides the most detailed spatial understanding of annotator reliability by operating on the full-resolution features from the decoder. This approach allows for fine-grained reliability estimation at the pixel level, capturing local variations in annotator reliability.

The pixel-wise reliability branch is implemented by applying three convolutional layers with batch normalization to the decoder's output features:

- Two initial convolutional layers with 32 filters each to process the features
- A final convolutional layer that produces  $R$  reliability maps, one for each annotator

This approach provides the most detailed spatial reliability information, allowing the model to capture pixel-level variations in annotator reliability. The resulting reliability maps have the same spatial dimensions as the input image, providing the finest granularity of reliability estimation among the three approaches.

## 5.3 Integration with TGCE Loss

The model produces two distinct outputs: segmentation masks  $\mathbf{\hat{Y}} = f(\mathbf{X}; \theta)$  and reliability maps  $\{\Lambda_r(\mathbf{X}; \theta)\}_{r=1}^R$ . These outputs work in tandem with the  $TGCE_{SS}$  loss function described in Section 4.1. The loss function simultaneously guides the learning of both the segmentation masks and reliability maps, ensuring that the model learns to balance the contributions of different annotators based on their estimated reliability.

The  $TGCE_{SS}$  loss function is implemented differently for each reliability approach, with specific adaptations to handle the different spatial dimensions of the reliability maps. Naturally, the model itself contains tunable hyperparameters related to the Generalized Cross Entropy (GCE) - Mean Absolute Error (MAE) trade off behavior (also known as  $q$  parameter) and the unreliability tolerance  $\tilde{C}$ , however, some additional hyperparameters were added with the purpose of regularizing the overall loss for every experiment.

### 1674 **5.3.1 Regularization terms**

1675 All three implementations include regularization terms to ensure proper learning of  
1676 the reliability maps:

#### 1677 **Numerical stability**

1678 To ensure numerical stability in the loss function computations, particularly in  
1679 division operations, a small epsilon value  $\epsilon$  is added to denominators. This  
1680 prevents division by zero errors and stabilizes gradient calculations during  
1681 training. The epsilon value is kept small enough to not significantly impact the loss  
1682 function behavior while providing the necessary numerical safeguards.

#### 1683 **Penalty for extreme reliabilities**

1684 Some optimization processes might favor a convergence to extreme reliability  
1685 estimations, by assigning either 0 or 1 to all labelers. To avoid this, a regularization  
1686 term is added to the loss function which penalizes the model for assigning extreme  
1687 reliability values and favoring those around a centroid. High scaling values for this  
1688 regularization term might lead to a convergence to a very homogeneous reliability  
1689 map across all labelers, which is not desired either.

$$\mathcal{L}_{centroid} = \alpha_{centroid} \cdot \mathbb{E}[(\Lambda_r - 0.5)^2] \quad (5-1)$$

#### 1690 **Entropy regularization**

1691 In this case, entropy is useful to ensure that the reliability maps are not too  
1692 homogeneous. The entropy regularization term encourages the model to assign  
1693 more discriminative reliability values by penalizing uniform distributions. This  
1694 prevents the model from converging to uninformative reliability estimates where  
1695 all labelers appear equally reliable, thus maintaining meaningful distinctions  
1696 between more and less trustworthy annotators.

$$\mathcal{L}_{entropy} = \alpha_{entropy} \cdot \mathbb{E}[\Lambda_r \log(1 + \Lambda_r) + (1 - \Lambda_r) \log(1 + (1 - \Lambda_r))] \quad (5-2)$$

### Sum-to-one regularization

The reliability branch for scalar, pixel and features approaches, might or might not be forced to sum to one through a softmax operation<sup>1</sup>. For the cases in which this sum is not forced to be unitary the flexibility of this sum is manageable through a regularization term which penalizes the model for not summing to one. This was found useful in situations in which the reliability maps were not expected to be punctuated (e.g. for the scalar reliability) as  $1/R$ , but rather close to the unit for all of them (if all labelers were good). This regularization term was found particularly useful in situations where labelers have good and similar levels of expertise.

$$\mathcal{L}_{sum} = \alpha_{sum} \cdot \mathbb{E}[\sum_{r \in R} (\Lambda_r - 1)^2] \quad (5-3)$$

### Missing labels penalizing

In situations where the datasets has a low labeling rate, the model might take an optimization shortcut by assigning high reliability values to labelers which did not label at all, since the element-wise product with the mask would be close to zero, which favors a minimization target, which is not desired. To avoid this, a term was added to the loss function which penalizes the model for assigning high reliability values to unlabeled patches, as follows:

$$\mathcal{L}_{\lambda} = \gamma \max \left( 0, \sum_{\forall r \in R} \frac{1}{1 - \alpha_{\lambda}} (\Lambda_r - \mu_r) \cdot e^{\frac{\Lambda_r - \mu_r - 1}{\beta}} \right) \quad (5-4)$$

Where  $\Lambda_r$  is the estimated reliability of the labeler  $r$  for a certain patch,  $\mu_r$  is the one hot label mask of the labeler  $r$  for the same patch,  $\alpha_{\lambda}$  and  $\beta$  are hyperparameters which control the penalty strength for high and low estimated

---

<sup>1</sup>This was left as flexible and configurable in the Network Architecture Building API from `seg_tgce`

1716 reliabilities, respectively and  $\gamma$  is a hyperparameter which controls the overall  
 1717 strength of the penalty.




Figure 5-3 Example of the  $\mathcal{L}_\lambda(z)$  function for values of  $\alpha_\lambda = 0.2$  and  $\beta = 1.0$ .

1718 This function applies no penalty for cases in which  $\Lambda_r < \mu_r$ , meaning no penalty  
 1719 for assigning any reliability value (ranging from 0 to 1) to a labeler who actually  
 1720 labeled a given patch. On the other hand, a penalty is indeed applied in cases where  
 1721  $\Lambda_r > \mu_r$ , meaning a penalty for assigning high reliability values to a labeler which  
 1722 did not label at all. The proportion on how high and how quickly (in the space of  
 1723 difference  $\Lambda_r - \mu_r$ ) the penalty is applied is controlled by the hyperparameters  $\alpha_\lambda$   
 1724 and  $\beta$ . The behavior of this term is shown in Figure 5-3.

1725 These regularization terms lead to the overall loss function being in the form:

$$\begin{aligned}
L_{TGCE_{reg}} = L_{TGCE_{ss}} & \quad (5-5) \\
& + \alpha_{reg} \cdot \mathbb{E}[(\Lambda_r - 0.5)^2] \\
& + \alpha_{entropy} \cdot \mathbb{E}[\Lambda_r \log(1 + \Lambda_r) + (1 - \Lambda_r) \log(1 + (1 - \Lambda_r))] \\
& + \alpha_{sum} \cdot \mathbb{E}[\sum_{\forall r \in R} (\Lambda_r - 1)^2] \\
& + \gamma \cdot \max \left( 0, \sum_{\forall r \in R} \frac{1}{1 - \alpha_\lambda} (\Lambda_r - \mu_r) \cdot e^{\frac{\Lambda_r - \mu_r - 1}{\beta}} \right)
\end{aligned}$$

For the scalar reliability approach, the loss function operates on global reliability values per annotator. The reliability maps are expanded to match the spatial dimensions of the predictions through tiling operations. This approach is computationally efficient but provides the same reliability value across all spatial locations.

The feature-based reliability approach requires additional processing to handle the spatial dimensions of the bottleneck features. The reliability maps are resized using bilinear interpolation to match the spatial dimensions of the predictions. This approach provides coarse-grained spatial reliability information while maintaining computational efficiency.

The pixel-wise reliability approach operates directly on full-resolution reliability maps, requiring no spatial transformations. This approach provides the most detailed spatial reliability information but requires more computational resources.

## 5.4 Training Process

The training process begins with the initialization of the ResNet-34 backbone using pre-trained weights, providing a strong foundation for feature extraction. The model's hyperparameters are optimized using Optuna [Akiba et al., 2019], a hyperparameter optimization framework that employs Bayesian optimization to search for optimal values of:

- Learning rate: Sampled logarithmically between  $10^{-5}$  and  $10^{-2}$
- TGCE<sub>SS</sub> parameters:
  - $q$ : Controls truncation of the loss function (0.1 to 0.9)
  - Noise tolerance: Controls model's robustness to label noise (0.1 to 0.9)
  - Regularization weights: Control the balance between different loss terms.

Two separate hyperparameter optimization studies were conducted using Optuna, each with 20 iterations, on the Oxford PET dataset with three scorers emulating SNR noises of -20, 0, and 10 dB respectively, and a labeling rate of 0.5. The first study focused exclusively on regularization parameters with preselected values of  $q = 0.66$ ,  $\tilde{C} = 0.62$ , and learning rate =  $6.5 \times 10^{-4}$ , aiming to determine the relative importance of different regularizers. The second study was comprehensive, including the mandatory loss function parameters ( $q$ ,  $\tilde{C}$ ) and learning rate, which are more sensitive to changes and directly affect the loss function and optimizer behavior.

Table 5-1 presents the hyperparameter importance ranking for the regularizers-only study, while Table 5-2 shows the results from the comprehensive study including all parameters.

| Hyperparameter             | Importance Score |
|----------------------------|------------------|
| $\beta_\lambda$            | 0.4728           |
| $\alpha_{\text{sum}}$      | 0.1364           |
| $\alpha_{\text{entropy}}$  | 0.1291           |
| $\alpha_\lambda$           | 0.1280           |
| $\alpha_{\text{centroid}}$ | 0.0997           |
| $\gamma_\lambda$           | 0.0340           |

Table 5-1 Hyperparameter importance ranking for regularization parameters obtained from Optuna optimization study after 20 iterations. This study used preselected values of  $q = 0.66$ ,  $\tilde{C} = 0.62$ , and learning rate =  $6.5 \times 10^{-4}$ . The importance scores indicate the relative impact of each regularization hyperparameter on the model's performance.

| Hyperparameter             | Importance Score |
|----------------------------|------------------|
| learning_rate              | 0.2267           |
| $q$                        | 0.1832           |
| $\tilde{C}$                | 0.1480           |
| $\alpha_\lambda$           | 0.1127           |
| $\alpha_{\text{sum}}$      | 0.0988           |
| $\alpha_{\text{entropy}}$  | 0.0868           |
| $\alpha_{\text{centroid}}$ | 0.0380           |
| $\beta_\lambda$            | 0.0821           |
| $\gamma_\lambda$           | 0.0237           |

Table 5-2 Comprehensive hyperparameter importance ranking including all parameters obtained from Optuna optimization study after 20 iterations. This study shows the relative importance of mandatory loss function parameters ( $q$ ,  $\tilde{C}$ ), learning rate, and regularization parameters together.

1764 After finding the optimal hyperparameters, the model is trained end-to-end using  
 1765 the Adam optimizer with the best learning rate found <sup>2</sup>. The TGCE<sub>SS</sub> loss function  
 1766 plays a crucial role in updating both the segmentation and reliability branches,  
 1767 ensuring that the model learns to effectively handle multiple annotators' inputs  
 1768 while accounting for their varying reliability.

1769 The model's architecture is specifically designed to address the challenges of  
 1770 multi-annotator segmentation. Through the ResNet-34 backbone, it learns robust  
 1771 segmentation features that capture high-level patterns in the data. The UNET's  
 1772 skip connections enable the preservation of fine-grained details, while the parallel  
 1773 reliability branch allows the model to adapt to annotator-specific characteristics.  
 1774 This comprehensive design enables the model to effectively handle multiple  
 1775 annotators' inputs while maintaining high segmentation accuracy and reliability  
 1776 estimation.

---

<sup>2</sup>For every experiment this was also tuned with the help of the Optuna study, having different optimal values for every experiment.




Figure 5-4 Training process of the proposed model overview. The model first undergoes hyperparameter optimization using Bayesian optimization, followed by end-to-end training with the optimal parameters. Estimated segmentation and reliability maps are computed for each input image, and the loss function is computed for the entire batch. Network parameters are tuned based on the loss function as objective (with minimization target) whereas the hyperparameters  $\phi$  tuning are based on the Dice coefficient of the validation data set (with maximization target).

## 5.5 Experimental Setup

### 5.5.1 Noisy Annotations Generation

A challenge arises for the creation of emulated noisy annotations datasets, since it is expected for images annotations to have some degree of expertise variability, similar to what is expected in real crowdsourced datasets. Simply introducing random noise into the annotated masks does not work, since the original morphological structures from the expected ground truth are far from being preserved. Instead, a “noisy” labeler is expected to produce an annotation which has at least some degree of morphological consistency on itself, even if it shows discrepancies when compared with some metric (like DICE score) against the ground truth annotation.



1788 This has been proven experimentally when introducing random noise into any  
 1789 popular segmentation dataset, in which the resulting mask is just a non coherent  
 1790 map of noise across the image, as shown in Figure 5-5.




Figure 5-5 Example of a noisy mask generated by naively introducing random noise into a ground truth mask. Morphological consistency is lost.

1791 For this reason, a more sophisticated approach was needed to create datasets with  
 1792 emulated noisy annotations. The approach used here was to use a pre-trained  
 1793 U-Net model to produce a good enough <sup>3</sup> segmentation of the image, which was  
 1794 then used to produce a “noisy” annotation by introducing random noise into the  
 1795 last encoder layers of the model, thus preserving the morphological consistency of  
 1796 the original annotation. This strategy works since slight modifications introduced  
 1797 into the encoder layers weights, resemble conceptual disturbances with respect to  
 1798 the original ground truth label, which kind of emulates the cultural behavior of  
 1799 having a different interpretation of the image, either for having a different level of  
 1800 expertise or for having a different point of view or school of thought. The first  
 1801 encoding layers were not modified, since it is expected human labelers would  
 1802 agree on the most fundamental structures (analog to extracted features from initial  
 1803 convolutional layers) in a similar way, thus preserving some degree of  
 1804 morphological consistency of the original annotation.

1805 In this way, the level of “disturbance” with respect to the ground truth for an

---

<sup>3</sup>More precisely, a U-Net architecture with Resnet-34 backbone model with a simple categorical cross-entropy loss which reached 0.8426 Dice Coefficient after 44 epochs.

emulated annotator can be controlled by the level of noise introduced into the weights of the last encoder layers, thus:

$$\mathbf{W}_{noisy} = \mathbf{W}_{original} + \epsilon_d \quad (5-6)$$

where  $\mathbf{W}_{original}$  represents the original weights of the last encoder layers,  $\epsilon_d$  follows a Gaussian distribution with zero mean and variance  $\sigma^2$ . The variance  $\sigma^2$  controls the level of noise introduced, and thus the degree of disturbance in the resulting segmentation masks.

With the application of the encoder layer weight perturbation technique, the resulting noisy masks are shown in Figure 5-6. It can be seen that the morphological consistency is preserved, even though the resulting masks are far from the ground truth annotations, which goes perfectly well for testing the robustness of the models against noisy annotations in crowdsourcing-like scenarios.




Figure 5-6 Noisy mask generated by enhancing the disturbances in the encoder layers weights for the Oxford-IIIT Pet Dataset. Morphological consistency is preserved. From left to right, SNR levels of noise in the encoder layer are 10, 5, 2, 0, -5 dB. Not adding noisy to the encoder layer would reach to an annotator behavior as good as the originally trained base model.

### 1818 5.5.2 Missing Labels Emulation

1819 It has also been determined that most real world crowdsourced datasets suffer from  
1820 missing labels, since not all annotators are able to label all of the input samples. To  
1821 emulate this behaviour, a missing labels dataset was created by randomly skipping  
1822 labels from the annotators labels based on a labeling rate  $r_l \in [0, 1]$ . The labeling rate  
1823  $r_l$  controls the proportion of labels that are skipped from the annotators labels. The  
1824 lower the labeling rate, the more labels are skipped from the annotators labels. The  
1825 implementation of this missing labels emulation together with the noisy annotators  
1826 generation can be flexibilized through an algorithm, which enables it to be used in  
1827 experiments with variable amount of labelers, noise levels, and labeling rates. An  
1828 overview of the algorithm is shown in Algorithm 5-1. A sample of the generated  
1829 noisy and missing labels dataset is shown in Figure 5-7.




Figure 5-7 Example of missing labels and noisy annotations emulation for the Oxford-IIIT Pet Dataset. The same label rate ( $r_l = 0.7$ ) is preserved for all labelers, which means in this case that every labeler is expected to label 70% of the samples, and skip the remaining 30%.

## 1830 5.6 Results

1831 Initially, a baseline model (ground truth only) was trained with an early stopping  
 1832 mechanism, which stopped the training when the validation loss stopped decreasing  
 1833 for 10 epochs, reaching only 20 epochs of significant training. The results are shown  
 1834 in Figure 5-8.

---

**Algorithm 5-1** Multi-Annotator Dataset Generation with Noisy and Missing Labels
 

---

**Require:** Base segmentation model  $M$ , SNR levels  $\{s_1, \dots, s_n\}$ , labeling rate  $r \in [0, 1]$

**Ensure:** Dataset with multiple noisy annotations and missing labels

```

1: Step 1: Generate Noisy Annotators
2: for each SNR level  $s_i$  do
3:   Create disturbed model  $M_i$  by adding noise to  $M$  with SNR  $s_i$ 
4: end for
5: Step 2: Initialize Labeler Assignment System
6: Create assignment matrix  $A \in \{0, 1\}^{N \times n}$  where  $N$  is number of samples
7: for each annotator  $j \in \{1, \dots, n\}$  do
8:   for each sample  $i \in \{1, \dots, N\}$  do
9:      $A_{i,j} \leftarrow \begin{cases} 1 & \text{with probability } r \\ 0 & \text{with probability } 1 - r \end{cases}$ 
10:   end for
11: end for
12: Step 3: Generate Multi-Annotator Dataset
13: for each image  $x$  in dataset do
14:   Resize  $x$  to target dimensions
15:   for each annotator model  $M_i$  do
16:     Generate annotation  $y_i = M_i(x)$ 
17:     Apply labeler assignment mask:  $y_i \leftarrow y_i \cdot A_{i,j}$ 
18:   end for
19:   Store  $(x, \{y_1, \dots, y_n\})$  in dataset
20: end for
   return Multi-annotator dataset with noisy and missing labels

```

---

Then, using the scalar model described in Section 5.2.1, training was performed for 20 epochs and Bayesian hyperparameter optimization on learning rate, decoder dropout rate, for the regularizers described in Section 5.3.1. The results are shown in Figure 5-9.




Figure 5-8 Segmentation results for baseline model on the Triple-Negative Breast Cancer (TNBC) (Crowd-Seg) dataset after 20 epochs of training. Even though the segmentation algorithm seems to show progress on its way to capture the underlying target structures to be segmented, since the ground truth annotation is not expressive enough, the Dice coefficient remains below 0.55, which is not considered a good performance. The baseline model remained with overfitting even after combining techniques of: Bayesian hyperparameter optimization, data augmentation, learning rate scheduling, dataset shuffling, and early stopping.




Figure 5-9 Segmentation results for scalar model on the TNBC (Crowd-Seg) dataset after 20 epochs of training. Again, training metrics show a constant growth, but not validation metrics, which shows a combined issue of overfitting and unreliable ground truth.

### 5.6.1 Oxford-IIIT Pet Dataset

The experimental results on the Oxford-IIIT Pet Dataset demonstrate interesting insights into the performance of different approaches. While STAPLE achieves the highest Dice and Jaccard scores (0.8718 and 0.7917 respectively) for the dataset with no missing labels (label rate of 1.0), our proposed TGCE<sub>SS</sub> approaches show superior performance for label rates below 1.0 (Table 5-3). The pixel-wise variant, while slightly lower in performance metrics, provides more detailed spatial reliability information. It is of note that the key strength of TGCE<sub>SS</sub> lies in its ability to simultaneously learn both segmentation and reliability estimation, providing valuable insights into annotator behavior and confidence levels that STAPLE and MV cannot offer.

| Method       | Label Rate 1.0  |                 | Label Rate 0.7      |                     | Label Rate 0.3      |                     |
|--------------|-----------------|-----------------|---------------------|---------------------|---------------------|---------------------|
|              | Dice            | Jaccard         | Dice                | Jaccard             | Dice                | Jaccard             |
| TGCE scalar  | 0.8622±0.009    | 0.7791±0.012    | <b>0.8568±0.017</b> | <b>0.7726±0.024</b> | <b>0.8529±0.004</b> | <b>0.7677±0.007</b> |
| TGCE feature | 0.6995±0.044    | 0.5733±0.047    | 0.4766±0.026        | 0.3441±0.027        | 0.4766±0.026        | 0.3441±0.027        |
| TGCE pixel   | 0.8468±0.018    | 0.7589±0.025    | 0.8545±0.013        | 0.7695±0.017        | 0.8283±0.001        | 0.7359±0.002        |
| MV           | 0.8132±0        | 0.7222±0        | 0.7572±0            | 0.6554±0            | 0.5566±0            | 0.4687±0            |
| STAPLE       | <b>0.8718±0</b> | <b>0.7917±0</b> | 0.8439±0            | 0.7609±0            | 0.5734±0            | 0.5158±0            |

Table 5-3 Performance metrics (Dice and Jaccard) at different labeling rates. Best values per category are bolded.

The reliability estimation capabilities of our model are particularly noteworthy, as demonstrated in Figures 5-10 and 5-11. The model successfully captures the varying reliability levels of different annotators, with higher reliability assigned to annotators with lower noise levels (10 dB) and lower reliability to those with higher noise levels (-20 dB). This ability to quantify and visualize annotator reliability is a significant advantage over traditional approaches like STAPLE and MV, which focus solely on segmentation accuracy without providing insights into annotator behavior.




Figure 5-10 Example evaluation of the proposed model in its last trained epoch. An accurate segmentation output is delivered as well as scalar reliability estimations from the labelers masks. Disturbance signal-to-noise ratios are:  $[-20, 0, 10]$  dB from left to right. Best found hyperparameters: learning rate = 0.0011,  $q = 0.4$ , noise tolerance = 0.2,  $\alpha_{reg} = 0.26$ ,  $\alpha_{entropy} = 0.03$ ,  $\alpha_{sum} = 0.27$ .




Figure 5-11 Example evaluation of the proposed model in some given epoch (5/50) of training, which demonstrates a quick convergence of the model for correctly estimating the reliability of the labelers masks. An accurate segmentation output is delivered as well as pixel-wise reliability estimations from the labelers masks. Disturbance signal-to-noise ratios are:  $[-20, 0, 10]$  dB from left to right. Best found hyperparameters: learning rate = 0.0027,  $q = 0.9$ , noise tolerance = 0.2,  $\alpha_{reg} = 0.38$ ,  $\alpha_{entropy} = 0.02$ ,  $\alpha_{sum} = 0.42$ .



## 5.6.2 TNBC Dataset

As mentioned in Section 2.4.2, [Amgad et al., 2019] presented a TNBC dataset, which is crowdsourced and available as HistomicsTK metadata. Later, [López-Pérez et al., 2023] sampled 151 Whole Slide Imagings (WSIs), from which 161 Region of Interests (ROIs) were extracted and labeled by 20 medical students. The dataset was sampled as patches of size 512x512, resulting in 10173 training patches, 1264 validation patches and 399 testing patches. The segmentation masks for the patches are composed of the classes: Tumor, Stroma, Inflammation, Necrosis, and Others<sup>4</sup>. This subset of the TNBC is named Crowd-Seg [López-Pérez et al., 2023].

Raw PNG files for the images patches and masks were provided by the authors, in practice, in order to optimize the use of these data, a custom tensorflow dataset was built for facilitating batch memory, caching and prefetching. The custom dataset implementation handles the loading and preprocessing of both images and masks, including normalization of images to the [0,1] range and one-hot encoding of the segmentation masks. It also manages the multi-annotator structure by creating a labeler mask that indicates which annotators provided labels for each sample. The dataset supports data augmentation through random transformations (flips, rotations, brightness/contrast adjustments) and includes visualization utilities for debugging and analysis purposes. This custom dataset implementation is publicly available<sup>5 6</sup> within the PyPI package included in the overall solution of this work. (see Section 6.3).

### Notes on the Crowd-Seg Dataset

The multiple labelers version of the TNBC dataset (Crowd-Seg) contains a couple of peculiarities that are worth mentioning. First, the label rate (inverse percentage of

---

<sup>4</sup>In practice, some of this masks also contain an additional class labeled as 'ignore'.

<sup>5</sup>[https://github.com/blotero/seg\\_tgce](https://github.com/blotero/seg_tgce)

<sup>6</sup><https://seg-tgce.readthedocs.io/en/latest/>

labels that are missing) is not uniform across the labelers and, in general, quite low, ranging from 1 to 12% as seen in Figure 5-12.




Figure 5-12 Label rate distribution across the labelers for the TNBC (Crowd-Seg) dataset.

On the other hand, the classes distribution across all patches are not equal, being ‘Tumor’ (class 2) the most frequent class. Despite the classes not being equally distributed, the distribution of the classes across the labelers is fairly uniform, as shown in Figure 5-13.




Figure 5-13 Classes distribution across the labelers for the TNBC (Crowd-Seg) dataset. Most common classes are ‘Tumor’ (class 2), ‘Other’ (class 1) and ‘Stroma’ (class 3).

It was also found that the dice score of the labelers masks against the ground truth differed across the labelers and the classes, as shown in Figure 5-14. The target class (Tumor) ranged dice scores across the labelers from 0.5 to 0.9, other classes can get as low as 0.4 for in average for some labelers, meaning very high levels of disagreement.




Figure 5-14 Dice score of the labelers masks against the ground truth for the TNBC (Crowd-Seg) dataset.

Finally, it was found that the expert labels, which would thereby be considered as ground truth for metrics evaluation, were not always correct, as shown in Figure 5-15. These morphological inconsistencies make an objective metric for evaluating the performance of the proposed model difficult to define, since the ground truth is not always correct or expressive enough.




Figure 5-15 Example of ground truth with odd geometries (second patch from top to bottom), which are not always correct or compatible with the underlying structures of the patch itself.



Figure 5-16 Training curves for the baseline model on the TNBC (Crowd-Seg) dataset.

## 5.7 Summary

The experimental results on both the Oxford-IIIT Pet and TNBC (Crowd-Seg) datasets provide valuable insights into the performance and limitations of our

1902 proposed approach. The model demonstrated exceptional performance on the  
1903 Oxford-IIIT Pet dataset, achieving competitive Dice and Jaccard scores while  
1904 successfully capturing annotator reliability patterns. The model's ability to  
1905 simultaneously learn segmentation and reliability estimation proved particularly  
1906 valuable, providing insights into annotator behavior that traditional approaches  
1907 like STAPLE and MV cannot offer.

1908 However, the experiments on the **TNBC** dataset revealed significant challenges in  
1909 evaluating semantic segmentation models for histopathological images when  
1910 dealing with unreliable ground truth. The morphological inconsistencies in the  
1911 expert annotations, combined with the extremely low label rates (1-12%) across  
1912 annotators, made it difficult to objectively measure the model's performance.  
1913 While the model showed promising convergence in training, the validation metrics  
1914 remained unstable, suggesting that the current dataset may not be suitable for  
1915 rigorous evaluation of multi-annotator segmentation approaches.

1916 These findings highlight the need for more reliable ground truth annotations in  
1917 histopathological image segmentation datasets. Future work should focus on:

- 1918 • Developing datasets with more consistent and reliable expert annotations
- 1919 • Implementing stricter quality control measures for ground truth generation
- 1920 • Exploring alternative evaluation metrics that account for annotation  
1921 uncertainty
- 1922 • Investigating methods to handle extremely sparse annotations more  
1923 effectively

1924 Despite these challenges, the model's performance on the Oxford-IIIT Pet dataset  
1925 demonstrates its effectiveness in handling multiple annotators with varying  
1926 reliability levels. The model successfully learned to:

- 1927 • Balance the contributions of different annotators based on their reliability

- 1928 • Provide detailed spatial reliability information through pixel-wise estimation
- 1929 • Maintain high segmentation accuracy even with missing labels
- 1930 • Adapt to different noise levels in annotator inputs

1931 These capabilities make the proposed approach particularly valuable for  
1932 real-world applications where multiple annotators with varying expertise levels  
1933 contribute to the dataset. The model's ability to simultaneously learn  
1934 segmentation and reliability estimation provides a more comprehensive  
1935 understanding of the annotation process, which could be leveraged to improve the  
1936 quality of future annotations and dataset construction.

1937

1938

1939

1940

---

CHAPTER

**SIX**

---

1941

FINAL REMARKS

1942

## 6.1 Conclusions

1943

1944

1945

1946

1947

1948

1949

1950

1951

This thesis has presented a novel approach to address the challenges of Image Semantic Segmentation (ISS) in histopathology images, particularly in crowdsourcing scenarios where multiple annotators with varying levels of expertise contribute to the labeling process. The proposed solution successfully tackles the fundamental challenges of modeling annotator performance and handling label inconsistencies while maintaining computational efficiency. The work demonstrates significant improvements over existing state-of-the-art methods, particularly in scenarios with noisy and missing labels, which are common in crowdsourcing approaches.

1952

1953

1954

1955

1956

1957

Throughout this work, several techniques have been explored to develop a model that can face Image Semantic Segmentation (ISS) of real world histopathology images in crowdsourcing scenarios. Different approaches were studied and evaluated to reach an optimal solution which could model the annotators' performance from the input space. The following remarks are based on the results obtained from the experiments conducted and reviewed literature:

- 1958 • Chained Gaussian Processes approaches, while robust for classification and  
1959 regression tasks, proved computationally intensive for image segmentation.  
1960 This limitation highlighted the need for more efficient approaches that could  
1961 handle the complexity of medical image segmentation while modeling  
1962 annotator performance.
- 1963 • The development of a novel loss function was crucial for the success of this  
1964 work, as existing loss functions were insufficient to model the annotators'  
1965 performance across the input space. The proposed loss function successfully  
1966 enabled the model to adaptively infer the best possible segmentation without  
1967 requiring separate inputs about labeler performance.
- 1968 • The introduction of a tensor map to codify labeler reliability proved to be an  
1969 effective approach, allowing the model to implicitly weigh the labelers based  
1970 on their performance across the mask and classes space, thus handling the  
1971 labeler performance inconsistency effectively.
- 1972 • The use of U-Shaped Deep Learning models as a building block for the  
1973 proposed solution was validated, as it effectively captured spatial  
1974 relationships between pixels and annotations while handling the variability  
1975 in image quality.
- 1976 • The combination of Chained Deep Learning approaches with the proposed  
1977 loss function and tensor map representation resulted in a comprehensive  
1978 solution being, to the best of our knowledge, the first approach on offering  
1979 an image segmentation solution for multiple labelers in medical image with  
1980 interpretability in the reliability of the labelers, successfully addressing both  
1981 labeler inconsistency and image quality variability. The model outperformed  
1982 other state of the art solutions in terms of segmentation accuracy in noisy  
1983 and missing labels conditions, which is rather common in crowdsourcing  
1984 approaches.



## 6.2 Future work

As tools and frameworks for bayesian approaches like Chained Gaussian Processes<sup>1</sup> evolve, it is expected to see a decrease in computational requirements and hence, making mixed models more appealing for the task of ISS of histopathology images. In that sense, the use of *Correlated Chained Gaussian Processes for Multiple Annotators (CCGPMA)* as a building block for the proposed solution could be a good approach, as it is a powerful model that can take advantage of the inter-dependencies between the annotators, and the powerfulness of CNNs to extract features from the images.

The noisy labelers emulation technique developed in this work presents an opportunity to test the model's robustness in various medical image segmentation challenges. This technique could be adapted to simulate different types of expert variability in other medical imaging domains:

- **Multi-modal Medical Imaging:** The technique could be extended to simulate expert disagreements in multi-modal scenarios, where different imaging modalities (e.g., MRI, CT, PET) might lead to varying interpretations of the same anatomical structure.
- **Temporal Consistency:** For longitudinal studies, the emulation could be modified to simulate how expert annotations might vary over time, helping to test the model's ability to maintain consistency in segmentation across different time points.
- **Cross-institutional Variability:** The noise injection technique could be calibrated to simulate the variations in annotation patterns that arise from different institutional protocols or training backgrounds.

---

<sup>1</sup>Including GPFlow, and GPFlux.

2009       • **Pathology-specific Challenges:** The emulation could be tailored to simulate  
2010       expert disagreements in specific pathological conditions, where the  
2011       morphological features might be more ambiguous or subject to  
2012       interpretation.

2013       These extensions would help validate the model's performance across a broader  
2014       range of medical imaging scenarios and ensure its robustness in real-world clinical  
2015       applications.

2016       The pixel-wise reliability mapping approach shows particular promise for future  
2017       development. While the current implementation already demonstrates good  
2018       performance, several enhancements could lead to significant improvements:

2019       • **Advanced Data Augmentation:** Implementing a comprehensive data  
2020       augmentation strategy could significantly improve the model's convergence  
2021       and performance. This could include:

- 2022               – Geometric transformations (rotations, flips, elastic deformations)
- 2023               – Intensity variations to simulate different staining conditions
- 2024               – Multi-scale training to handle varying tissue sizes
- 2025               – Adversarial training to improve robustness

2026       These augmentations would help the model learn more robust features and  
2027       better generalize to unseen data.

2028       • **Attention Mechanisms:** Integrating attention mechanisms into the pixel-wise  
2029       reliability branch could help the model focus on the most relevant regions for  
2030       reliability estimation, potentially improving both accuracy and computational  
2031       efficiency.

2032       • **Multi-scale Reliability Estimation:** Implementing a hierarchical approach to  
2033       reliability estimation could capture both fine-grained local variations and  
2034       broader regional patterns in annotator reliability, this would imply a  
2035       network with yet more branches and outputs.

The combination of these enhancements with the existing pixel-wise reliability mapping approach could potentially outperform all other techniques, as it would provide the most detailed and accurate representation of annotator reliability while maintaining computational efficiency. The use of data augmentation, in particular, could help the model converge faster and achieve better generalization, as it would expose the model to a wider variety of possible variations in both the input images and the corresponding annotations.

## 6.3 Academic Discussion

To facilitate the adoption and further development of the proposed approach, several resources have been made publicly available. A comprehensive GitHub repository<sup>2</sup> has been created that provides easy-to-use APIs for implementing the proposed solution. The repository includes pre-built dataset loaders for TensorFlow/Keras frameworks, implementation of the  $TGCE_{SS}$  loss function in all its variants (scalar, features, and pixel-wise), and a neural network architecture builder with experimental examples. The two `tfds` datasets available in the repository are:

- **Triple Negative Breast Cancer (TNBC) Dataset:** A dataset of histopathology images of breast cancer, with annotations of the tumor regions from [López-Pérez et al., 2024]. The dataset is automatically fetched from AWS S3, which eliminates the need for manual downloads or external dependencies.
- **Oxford-IIIT Pet Dataset:** A dataset of pet images, based on the well-known Oxford-IIIT Pet Dataset with annotations of the pet regions. The mechanism for emulating noisy crowdsourced annotations exposed throughout the present work (see Section 5.5.1) is also available in the repository in a tunable manner, just as the TNBC dataset, this `tfds` is batteries included

---

<sup>2</sup>[https://github.com/blotero/seg\\_tgce](https://github.com/blotero/seg_tgce)

2061 (automatic fetching and mapping, no need to manually download any  
2062 external source).

2063 To further simplify the adoption of the proposed solution, a PyPI package has been  
2064 published with detailed documentation available at  
2065 <https://seg-tgce.readthedocs.io/en/latest/>. The documentation provides  
2066 comprehensive installation instructions, API reference, usage examples, and best  
2067 practices for implementation. This package makes it straightforward for the  
2068 community to integrate the proposed approach into their existing workflows, with  
2069 a focus on ease of use and clear documentation of all available features and  
2070 options.

2071 An academic publication is currently being prepared for submission to a Q1  
2072 journal. This publication will provide detailed theoretical foundations of the  
2073 proposed approach, comprehensive experimental results and analysis, comparison  
2074 with state-of-the-art methods, and discussion of limitations and future directions.  
2075 The publication will serve as the primary reference for the theoretical and practical  
2076 aspects of the proposed solution, ensuring that the work is properly documented  
2077 and validated through the peer-review process.

2078 These resources collectively contribute to the reproducibility and accessibility of  
2079 the research, enabling the broader scientific community to build upon and extend  
2080 the proposed approach. The combination of open-source code, well-documented  
2081 APIs, and peer-reviewed publication ensures that the work can be properly  
2082 evaluated, validated, and utilized by other researchers in the field. This  
2083 comprehensive approach to sharing research outputs aligns with current best  
2084 practices in scientific computing and Machine Learning research, promoting both  
2085 transparency and practical utility of the proposed solution.

## BIBLIOGRAPHY

- 2087 [Abadi et al., 2016] Abadi, M., Barham, P., Chen, J., Chen, Z., Davis, A., Dean, J.,  
2088 Devin, M., Ghemawat, S., Irving, G., Isard, M., Kudlur, M., Levenberg, J., Monga,  
2089 R., Moore, S., Murray, D. G., Steiner, B., Tucker, P. A., Vasudevan, V., Warden,  
2090 P., Wicke, M., Yu, Y., and Zhang, X. (2016). Tensorflow: A system for large-scale  
2091 machine learning. CoRR, abs/1605.08695. (page 37)
- 2092 [Akiba et al., 2019] Akiba, T., Sano, S., Yanase, T., Ohta, T., and Koyama, M. (2019).  
2093 Optuna: A next-generation hyperparameter optimization framework. (page 75)
- 2094 [Alloghani et al., 2020] Alloghani, M., Al-Jumeily, D., Mustafina, J., Hussain, A., and  
2095 Aljaaf, A. J. (2020). *A Systematic Review on Supervised and Unsupervised Machine*  
2096 *Learning Algorithms for Data Science*, pages 3–21. Springer International Publishing,  
2097 Cham. (page 35)
- 2098 [Amgad et al., 2019] Amgad, M., Elfandy, H., Hussein, H., Atteya, L. A., Elsebaie,  
2099 M. A. T., Abo Elnasr, L. S., Sakr, R. A., Salem, H. S. E., Ismail, A. F., Saad, A. M.,  
2100 Ahmed, J., Elsebaie, M. A. T., Rahman, M., Ruhban, I. A., Elgazar, N. M., Alagha, Y.,  
2101 Osman, M. H., Alhusseiny, A. M., Khalaf, M. M., Younes, A.-A. F., Abdulkarim, A.,  
2102 Younes, D. M., Gadallah, A. M., Elkashash, A. M., Fala, S. Y., Zaki, B. M., Beezley,  
2103 J., Chittajallu, D. R., Manthey, D., Gutman, D. A., and Cooper, L. A. D. (2019).  
2104 Structured crowdsourcing enables convolutional segmentation of histology  
2105 images. *Bioinformatics*, 35(18):3461–3467. (pages xxii, 29, 40, 42, and 87)

- 2106 [Avanzo et al., 2024] Avanzo, M., Stancanella, J., Pirrone, G., Drigo, A., and Retico,  
2107 A. (2024). The evolution of artificial intelligence in medical imaging: From  
2108 computer science to machine and deep learning. *Cancers (Basel)*, 16(21):3702.  
2109 Author Joseph Stancanella is employed by Elekta SA. The remaining authors  
2110 declare no commercial or financial conflicts of interest. (page 3)
- 2111 [Azad et al., 2024] Azad, R., Aghdam, E. K., Rauland, A., Jia, Y., Avval, A. H.,  
2112 Bozorgpour, A., Karimijafarbigloo, S., Cohen, J. P., Adeli, E., and Merhof, D.  
2113 (2024). Medical image segmentation review: The success of u-net. *IEEE*  
2114 *Transactions on Pattern Analysis and Machine Intelligence*, 46(12):10076–10095.  
2115 (pages 2 and 15)
- 2116 [Banerjee et al., 2025] Banerjee, A., Shan, H., and Feng, R. (2025). Editorial:  
2117 Artificial intelligence applications for cancer diagnosis in radiology. *Frontiers in*  
2118 *Radiology*, 5. (page 9)
- 2119 [Bhalgat et al., 2018] Bhalgat, Y., Shah, M. P., and Awate, S. P. (2018). Annotation-  
2120 cost minimization for medical image segmentation using suggestive mixed  
2121 supervision fully convolutional networks. *CoRR*, abs/1812.11302. (page 5)
- 2122 [Bingham et al., 2020] Bingham, G., Macke, W., and Miikkulainen, R. (2020).  
2123 Evolutionary optimization of deep learning activation functions. In *Proceedings of*  
2124 *the 2020 Genetic and Evolutionary Computation Conference*, GECCO '20, page 289–296,  
2125 New York, NY, USA. Association for Computing Machinery. (page 36)
- 2126 [Brito-Pacheco et al., 2025] Brito-Pacheco, D., Giannopoulos, P., and Reyes-  
2127 Aldasoro, C. C. (2025). Persistent homology in medical image processing: A  
2128 literature review. (page 2)
- 2129 [Carmo et al., 2025] Carmo, D. S., Pezzulo, A. A., Villacreses, R. A., Eisenbeisz,  
2130 M. L., Anderson, R. L., Van Dorin, S. E., Rittner, L., Lotufo, R. A., Gerard, S. E.,  
2131 Reinhardt, J. M., and Comellas, A. P. (2025). Manual segmentation of opacities  
2132 and consolidations on ct of long covid patients from multiple annotators. *Scientific*  
2133 *Data*, 12(1):402. (page 9)

- [de G. Matthews et al., 2016] de G. Matthews, A. G., van der Wilk, M., Nickson, T., Fujii, K., Boukouvalas, A., León-Villagrà, P., Ghahramani, Z., and Hensman, J. (2016). Gpflow: A gaussian process library using tensorflow. (page 48)
- [Duarte and Baraniuk, 2012] Duarte, M. F. and Baraniuk, R. G. (2012). Kronecker compressive sensing. *IEEE Transactions on Image Processing*, 21(2):494–504. (page 27)
- [Dutordoir et al., 2021] Dutordoir, V., Salimbeni, H., Hambro, E., McLeod, J., Leibfried, F., Artemev, A., van der Wilk, M., Hensman, J., Deisenroth, M. P., and John, S. (2021). Gpflux: A library for deep gaussian processes. (page 49)
- [Elhaminia et al., 2025] Elhaminia, B., Alsalemi, A., Nasir, E., Jahanifar, M., Awan, R., Young, L. S., Rajpoot, N. M., Minhas, F., and Raza, S. E. A. (2025). From traditional to deep learning approaches in whole slide image registration: A methodological review. (page 2)
- [Elnakib et al., 2020] Elnakib, A., Elmenabawy, N., and S Moustafa, H. (2020). Automated deep system for joint liver and tumor segmentation using majority voting. *MEJ-Mansoura Engineering Journal*, 45(4):30–36. (page 12)
- [Farahani et al., 2015] Farahani, N., Parwani, A. V., and Pantanowitz, L. (2015). Whole slide imaging in pathology: advantages, limitations, and emerging perspectives. *Pathology and Laboratory Medicine International*, 7:23–33. (pages xxi and 30)
- [Ghezloo et al., 2025] Ghezloo, F., Chang, O. H., Knezevich, S. R., Shaw, K. C., Thigpen, K. G., Reisch, L. M., Shapiro, L. G., and Elmore, J. G. (2025). Robust roi detection in whole slide images guided by pathologists’ viewing patterns. *Journal of Imaging Informatics in Medicine*, 38(1):439–454. (pages xxi, 31, and 32)
- [Gil-Gonzalez et al., 2018] Gil-Gonzalez, J., Alvarez-Meza, A., and Orozco-Gutierrez, A. (2018). Learning from multiple annotators using kernel alignment. *Pattern Recognition Letters*, 116:150–156. (pages 43 and 57)

- [Gil-Gonzalez et al., 2023] Gil-Gonzalez, J., Giraldo, J. J., Alvarez-Meza, A. M., Orozco-Gutierrez, A., and Alvarez, M. A. (2023). Correlated chained gaussian processes for datasets with multiple annotators. *IEEE Transactions on Neural Networks and Learning Systems*, 34(8):4514–4528. (pages 22, 50, and 59)
- [Gil-Gonzalez et al., 2021a] Gil-Gonzalez, J., Orozco-Gutierrez, A., and Alvarez-Meza, A. (2021a). Learning from multiple inconsistent and dependent annotators to support classification tasks. *Neurocomputing*, 423:236–247. (pages 44, 45, and 57)
- [Gil-Gonzalez et al., 2021b] Gil-Gonzalez, J., Valencia-Duque, A., Alvarez-Meza, A., Orozco-Gutierrez, A., and Garcia-Moreno, A. (2021b). Regularized chained deep neural network classifier for multiple annotators. *Applied Sciences*, 11(12). (pages 46, 47, and 57)
- [Gil-González and Álvarez Meza, 2023] Gil-González, J. and Álvarez Meza, A. M. (2023). A supervised learning framework in the context of multiple annotators. (pages 18, 50, 57, and 58)
- [Giraldo et al., 2022] Giraldo, J.-J., Zhang, J., and Álvarez, M. A. (2022). Correlated chained gaussian processes for modelling citizens mobility using a zero-inflated poisson likelihood. *IEEE Transactions on Intelligent Transportation Systems*, 23(11):20337–20351. (pages 45 and 46)
- [Giri and Bhatia, 2024] Giri, K. and Bhatia, S. (2024). Artificial intelligence in nephrology- its applications from bench to bedside. *International Journal of Advances in Nephrology Research*, 7(1):90–97. (page 8)
- [Grefve et al., 2024] Grefve, J., Söderkvist, K., Gunnlaugsson, A., Sandgren, K., Jonsson, J., Keeratijarut Lindberg, A., Nilsson, E., Axelsson, J., Bergh, A., Zackrisson, B., Moreau, M., Thellenberg Karlsson, C., Olsson, L., Widmark, A., Riklund, K., Blomqvist, L., Berg Loegager, V., Strandberg, S. N., and Nyholm, T. (2024). Histopathology-validated gross tumor volume delineations of intraprostatic lesions using psma-positron emission tomography/multiparametric magnetic resonance imaging. *Physics and Imaging in Radiation Oncology*, 31:100633. (page 14)



- [Habis, 2024] Habis, A. A. (2024). *Developing interactive artificial intelligence tools to assist pathologists with histology annotation*. Theses, Institut Polytechnique de Paris. (page 8)
- [He et al., 2015] He, K., Zhang, X., Ren, S., and Sun, J. (2015). Deep residual learning for image recognition. (page 68)
- [Hossain et al., 2023] Hossain, M. S., Shahriar, G. M., Syeed, M. M. M., Uddin, M. F., Hasan, M., Shivam, S., and Advani, S. (2023). Region of interest (roi) selection using vision transformer for automatic analysis using whole slide images. *Scientific Reports*, 13(1):11314. (pages 31 and 32)
- [Hu et al., 2025] Hu, D., Jiang, Z., Shi, J., Xie, F., Wu, K., Tang, K., Cao, M., Huai, J., and Zheng, Y. (2025). Pathology report generation from whole slide images with knowledge retrieval and multi-level regional feature selection. *Computer Methods and Programs in Biomedicine*, 263:108677. (pages 2 and 30)
- [Jia et al., 2014] Jia, Y., Shelhamer, E., Donahue, J., Karayev, S., Long, J., Girshick, R. B., Guadarrama, S., and Darrell, T. (2014). Caffe: Convolutional architecture for fast feature embedding. *CoRR*, abs/1408.5093. (page 37)
- [Karthikeyan et al., 2023] Karthikeyan, R., McDonald, A., and Mehta, R. (2023). What’s in a label? annotation differences in forecasting mental fatigue using ecg data and seq2seq architectures. (page 10)
- [Kendall and Gal, 2017] Kendall, A. and Gal, Y. (2017). What uncertainties do we need in bayesian deep learning for computer vision? In Guyon, I., Luxburg, U. V., Bengio, S., Wallach, H., Fergus, R., Vishwanathan, S., and Garnett, R., editors, *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc. (page 60)
- [Kim et al., 2024] Kim, Y., Lee, E., Lee, Y., and Oh, U. (2024). Understanding novice’s annotation process for 3d semantic segmentation task with human-in-the-loop. In *Proceedings of the 29th International Conference on Intelligent User Interfaces*, IUI ’24, page 444–454, New York, NY, USA. Association for Computing Machinery. (page 10)

- 2220 [Lam and Suen, 1997] Lam, L. and Suen, S. (1997). Application of majority  
2221 voting to pattern recognition: an analysis of its behavior and performance.  
2222 *IEEE Transactions on Systems, Man, and Cybernetics - Part A: Systems and Humans*,  
2223 27(5):553–568. (pages 12 and 22)
- 2224 [Lin et al., 2024] Lin, Y., Lian, A., Liao, M., and Yuan, S. (2024). Bcdnet: A fast  
2225 residual neural network for invasive ductal carcinoma detection. (page 8)
- 2226 [Liu et al., 2020] Liu, H., Ong, Y.-S., Shen, X., and Cai, J. (2020). When gaussian  
2227 process meets big data: A review of scalable gps. *IEEE Transactions on Neural  
2228 Networks and Learning Systems*, 31(11):4405–4423. (pages 46 and 60)
- 2229 [López-Pérez et al., 2023] López-Pérez, M., Morales-Álvarez, P., Cooper, L. A. D.,  
2230 Molina, R., and Katsaggelos, A. K. (2023). Crowdsourcing segmentation  
2231 of histopathological images using annotations provided by medical students.  
2232 In Juarez, J. M., Marcos, M., Stiglic, G., and Tucker, A., editors, *Artificial  
2233 Intelligence in Medicine*, pages 245–249, Cham. Springer Nature Switzerland.  
2234 (pages 5, 6, 9, and 87)
- 2235 [Lu et al., 2023] Lu, X., Ratcliffe, D., Kao, T.-T., Tikhonov, A., Litchfield, L., Rodger,  
2236 C., and Wang, K. (2023). Rethinking quality assurance for crowdsourced multi-  
2237 roi image segmentation. *Proceedings of the AAAI Conference on Human Computation  
2238 and Crowdsourcing*, 11(1):103–114. (pages 5 and 9)
- 2239 [López-Pérez et al., 2024] López-Pérez, M., Morales-Álvarez, P., Cooper, L. A.,  
2240 Felicelli, C., Goldstein, J., Vadasz, B., Molina, R., and Katsaggelos, A. K.  
2241 (2024). Learning from crowds for automated histopathological image  
2242 segmentation. *Computerized Medical Imaging and Graphics*, 112:102327.  
2243 (pages xxi, 5, 16, 18, and 97)
- 2244 [Mazzarini et al., 2021] Mazzarini, M., Falchi, M., Bani, D., and Migliaccio, A. R.  
2245 (2021). Evolution and new frontiers of histology in bio-medical research.  
2246 *Microscopy Research and Technique*, 84(2):217–237. (pages xxi and 29)

- 2247 [Pan et al., 2021] Pan, X., Lu, Y., Lan, R., Liu, Z., Qin, Z., Wang, H., and Liu, Z. (2021).  
2248 Mitosis detection techniques in h&e stained breast cancer pathological images:  
2249 A comprehensive review. *Computers & Electrical Engineering*, 91:107038. (page 33)
- 2250 [Panayides et al., 2020] Panayides, A. S., Amini, A., Filipovic, N. D., Sharma, A.,  
2251 Tsiftaris, S. A., Young, A., Foran, D., Do, N., Golemati, S., Kurc, T., Huang, K.,  
2252 Nikita, K. S., Veasey, B. P., Zervakis, M., Saltz, J. H., and Pattichis, C. S. (2020). Ai  
2253 in medical imaging informatics: Current challenges and future directions. *IEEE*  
2254 *Journal of Biomedical and Health Informatics*, 24(7):1837–1857. (page 2)
- 2255 [Parkhi et al., 2012] Parkhi, O. M., Vedaldi, A., Zisserman, A., and Jawahar, C. V.  
2256 (2012). Cats and dogs. In *IEEE Conference on Computer Vision and Pattern Recognition*.  
2257 (page 38)
- 2258 [Paszke et al., 2019] Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan,  
2259 G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., Desmaison, A., Köpf, A., Yang,  
2260 E. Z., DeVito, Z., Raison, M., Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., Bai,  
2261 J., and Chintala, S. (2019). Pytorch: An imperative style, high-performance deep  
2262 learning library. *CoRR*, abs/1912.01703. (page 37)
- 2263 [Plasilova et al., 2016] Plasilova, M. L., Hayse, B., Killelea, B. K., Horowitz, N. R.,  
2264 Chagpar, A. B., and Lannin, D. R. (2016). Features of triple-negative breast cancer:  
2265 Analysis of 38,813 cases from the national cancer database. *Medicine (Baltimore)*,  
2266 95(35):e4614. (page 41)
- 2267 [Qiu et al., 2022] Qiu, Y., Hu, Y., Kong, P., Xie, H., Zhang, X., Cao, J., Wang, T.,  
2268 and Lei, B. (2022). Automatic prostate gleason grading using pyramid semantic  
2269 parsing network in digital histopathology. *Frontiers in Oncology*, 12. (page 14)
- 2270 [Rashmi et al., 2021] Rashmi, R., Prasad, K., and Udupa, C. B. K. (2021).  
2271 Breast histopathological image analysis using image processing techniques for  
2272 diagnostic purposes: A methodological review. *Journal of Medical Systems*, 46(1):7.  
2273 (pages 1 and 5)

- 2274 [Rasmussen and Williams, 2006a] Rasmussen, C. E. and Williams, C. K. I. (2006a).  
 2275 *Gaussian Processes for Machine Learning*. MIT Press, Cambridge, MA, USA. (page 46)
- 2276 [Rasmussen and Williams, 2006b] Rasmussen, C. E. and Williams, C. K. I. (2006b).  
 2277 *Gaussian Processes for Machine Learning*. MIT Press, Cambridge, MA, USA.  
 2278 (pages 46, 58, and 59)
- 2279 [Roelofs et al., 2019] Roelofs, R., Shankar, V., Recht, B., Fridovich-Keil, S., Hardt,  
 2280 M., Miller, J., and Schmidt, L. (2019). A meta-analysis of overfitting in machine  
 2281 learning. In Wallach, H., Larochelle, H., Beygelzimer, A., d'Alché-Buc, F.,  
 2282 Fox, E., and Garnett, R., editors, *Advances in Neural Information Processing Systems*,  
 2283 volume 32. Curran Associates, Inc. (page 36)
- 2284 [Ronneberger et al., 2015] Ronneberger, O., Fischer, P., and Brox, T. (2015). U-  
 2285 net: Convolutional networks for biomedical image segmentation. In Navab, N.,  
 2286 Hornegger, J., Wells, W. M., and Frangi, A. F., editors, *Medical Image Computing  
 2287 and Computer-Assisted Intervention – MICCAI 2015*, pages 234–241, Cham. Springer  
 2288 International Publishing. (pages xxii, 15, 22, and 68)
- 2289 [Ryou et al., 2025] Ryou, H., Thomas, E., Wojciechowska, M., Harding, L., Tam,  
 2290 K. H., Wang, R., Hu, X., Rittscher, J., Cooper, R., and Royston, D. (2025). Reticulin-  
 2291 free quantitation of bone marrow fibrosis in mpns: Utility and applications.  
 2292 *eJHaem*, 6(2):e70005. (page 2)
- 2293 [Sarvamangala and Kulkarni, 2022] Sarvamangala, D. R. and Kulkarni, R. V. (2022).  
 2294 Convolutional neural networks in medical image understanding: a survey.  
 2295 *Evolutionary Intelligence*, 15(1):1–22. (pages 3 and 8)
- 2296 [Saul et al., 2016] Saul, A. D., Hensman, J., Vehtari, A., and Lawrence, N. D. (2016).  
 2297 Chained gaussian processes. (pages 45 and 48)
- 2298 [Shah et al., 2018] Shah, M. P., Merchant, S. N., and Awate, S. P. (2018).  
 2299 Ms-net: Mixed-supervision fully-convolutional networks for full-resolution  
 2300 segmentation. In Frangi, A. F., Schnabel, J. A., Davatzikos, C., Alberola-  
 2301 López, C., and Fichtinger, G., editors, *Medical Image Computing and Computer  
 2302 Assisted Intervention – MICCAI 2018*, pages 379–387, Cham. Springer International  
 2303 Publishing. (page 5)

- 2304 [Shalf, 2020] Shalf, J. (2020). The future of computing beyond moore's law.  
2305 *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering*  
2306 *Sciences*, 378(2166):20190061. (page 3)
- 2307 [Soydaner, 2020] Soydaner, D. (2020). A comparison of optimization algorithms  
2308 for deep learning. *International Journal of Pattern Recognition and Artificial Intelligence*,  
2309 34(13):2052013. (page 36)
- 2310 [Taye, 2023] Taye, M. M. (2023). Theoretical understanding of convolutional neural  
2311 network: Concepts, architectures, applications, future directions. *Computation*,  
2312 11(3). (page 35)
- 2313 [TIAN and Zhu, 2015] TIAN, T. and Zhu, J. (2015). Max-margin majority voting for  
2314 learning from crowds. In Cortes, C., Lawrence, N., Lee, D., Sugiyama, M., and  
2315 Garnett, R., editors, *Advances in Neural Information Processing Systems*, volume 28.  
2316 Curran Associates, Inc. (page 13)
- 2317 [Triana-Martinez et al., 2023] Triana-Martinez, J. C., Gil-González, J., Fernandez-  
2318 Gallego, J. A., Álvarez Meza, A. M., and Castellanos-Dominguez, C. G. (2023).  
2319 Chained deep learning using generalized cross-entropy for multiple annotators  
2320 classification. *Sensors*, 23(7). (page 21)
- 2321 [Warfield et al., 2004] Warfield, S., Zou, K., and Wells, W. (2004). Simultaneous  
2322 truth and performance level estimation (staple): an algorithm for the validation  
2323 of image segmentation. *IEEE Transactions on Medical Imaging*, 23(7):903–921.  
2324 (pages 13 and 22)
- 2325 [Weitz et al., 2023] Weitz, P., Valkonen, M., Solorzano, L., Carr, C., Kartasalo, K.,  
2326 Boissin, C., Koivukoski, S., Kuusela, A., Rasic, D., Feng, Y., Sinius Pouplier,  
2327 S., Sharma, A., Ledesma Eriksson, K., Latonen, L., Laenkholm, A.-V., Hartman,  
2328 J., Ruusuvuori, P., and Rantalainen, M. (2023). A multi-stain breast cancer  
2329 histological whole-slide-image data set from routine diagnostics. *Scientific Data*,  
2330 10(1):562. (pages xxii, 34, 39, and 40)

- 2331 [Werbos, 1990] Werbos, P. (1990). Backpropagation through time: what it does  
2332 and how to do it. *Proceedings of the IEEE*, 78(10):1550–1560. (page 34)
- 2333 [Wilson et al., 2016] Wilson, A. G., Hu, Z., Salakhutdinov, R., and Xing, E. P. (2016).  
2334 Deep kernel learning. In Gretton, A. and Robert, C. C., editors, *Proceedings of*  
2335 *the 19th International Conference on Artificial Intelligence and Statistics*, volume 51  
2336 of *Proceedings of Machine Learning Research*, pages 370–378, Cadiz, Spain. PMLR.  
2337 (page 60)
- 2338 [Xu et al., 2024] Xu, Y., Quan, R., Xu, W., Huang, Y., Chen, X., and Liu, F. (2024).  
2339 Advances in medical image segmentation: A comprehensive review of traditional,  
2340 deep learning and hybrid approaches. *Bioengineering*, 11(10). (pages 3 and 8)
- 2341 [Yan et al., 2014] Yan, Y., Rosales, R., Fung, G., Subramanian, R., and Dy, J. (2014).  
2342 Learning from multiple annotators with varying expertise. *Machine Learning*,  
2343 95(3):291–327. (page 43)
- 2344 [Yu et al., 2025] Yu, J., Li, B., Pan, X., Shi, Z., Wang, H., Lan, R., and Luo, X. (2025).  
2345 Semi-supervised gland segmentation via feature-enhanced contrastive learning  
2346 and dual-consistency strategy. *IEEE Journal of Biomedical and Health Informatics*,  
2347 pages 1–11. (page 2)
- 2348 [Zhang, 2022] Zhang, J. (2022). Knowledge learning with crowdsourcing: A  
2349 brief review and systematic perspective. *IEEE/CAA Journal of Automatica Sinica*,  
2350 9(5):749–762. (page 43)
- 2351 [Zhang et al., 2020] Zhang, L., Tanno, R., Xu, M.-C., Jin, C., Jacob, J., Ciccarrelli, O.,  
2352 Barkhof, F., and Alexander, D. (2020). Disentangling human error from ground  
2353 truth in segmentation of medical images. In Larochelle, H., Ranzato, M., Hadsell,  
2354 R., Balcan, M., and Lin, H., editors, *Advances in Neural Information Processing Systems*,  
2355 volume 33, pages 15750–15762. Curran Associates, Inc. (page 17)
- 2356 [Zhang and Sabuncu, 2018] Zhang, Z. and Sabuncu, M. R. (2018). Generalized  
2357 cross entropy loss for training deep neural networks with noisy labels.  
2358 (pages 20, 22, and 62)

- 2359 [Zhao et al., 2020] Zhao, R., Qian, B., Zhang, X., Li, Y., Wei, R., Liu, Y., and Pan,  
2360 Y. (2020). Rethinking dice loss for medical image segmentation. In *2020 IEEE*  
2361 *International Conference on Data Mining (ICDM)*, pages 851–860. (page 20)
- 2362 [Zhou et al., 2021] Zhou, S. K., Greenspan, H., Davatzikos, C., Duncan, J. S.,  
2363 Van Ginneken, B., Madabhushi, A., Prince, J. L., Rueckert, D., and Summers, R. M.  
2364 (2021). A review of deep learning in medical imaging: Imaging traits, technology  
2365 trends, case studies with progress highlights, and future promises. *Proceedings of*  
2366 *the IEEE*, 109(5):820–838. (pages 1 and 2)
- 2367 [Zhou et al., 2024] Zhou, Z., Gong, H., Hsieh, S., McCollough, C. H., and Yu, L.  
2368 (2024). Image quality evaluation in deep-learning-based ct noise reduction using  
2369 virtual imaging trial methods: Contrast-dependent spatial resolution. *Medical*  
2370 *Physics*, 51(8):5399–5413. (page 10)